

Tutorial 8: Bilateral Filtering Non-Local Means

Digital Image Processing (236860)



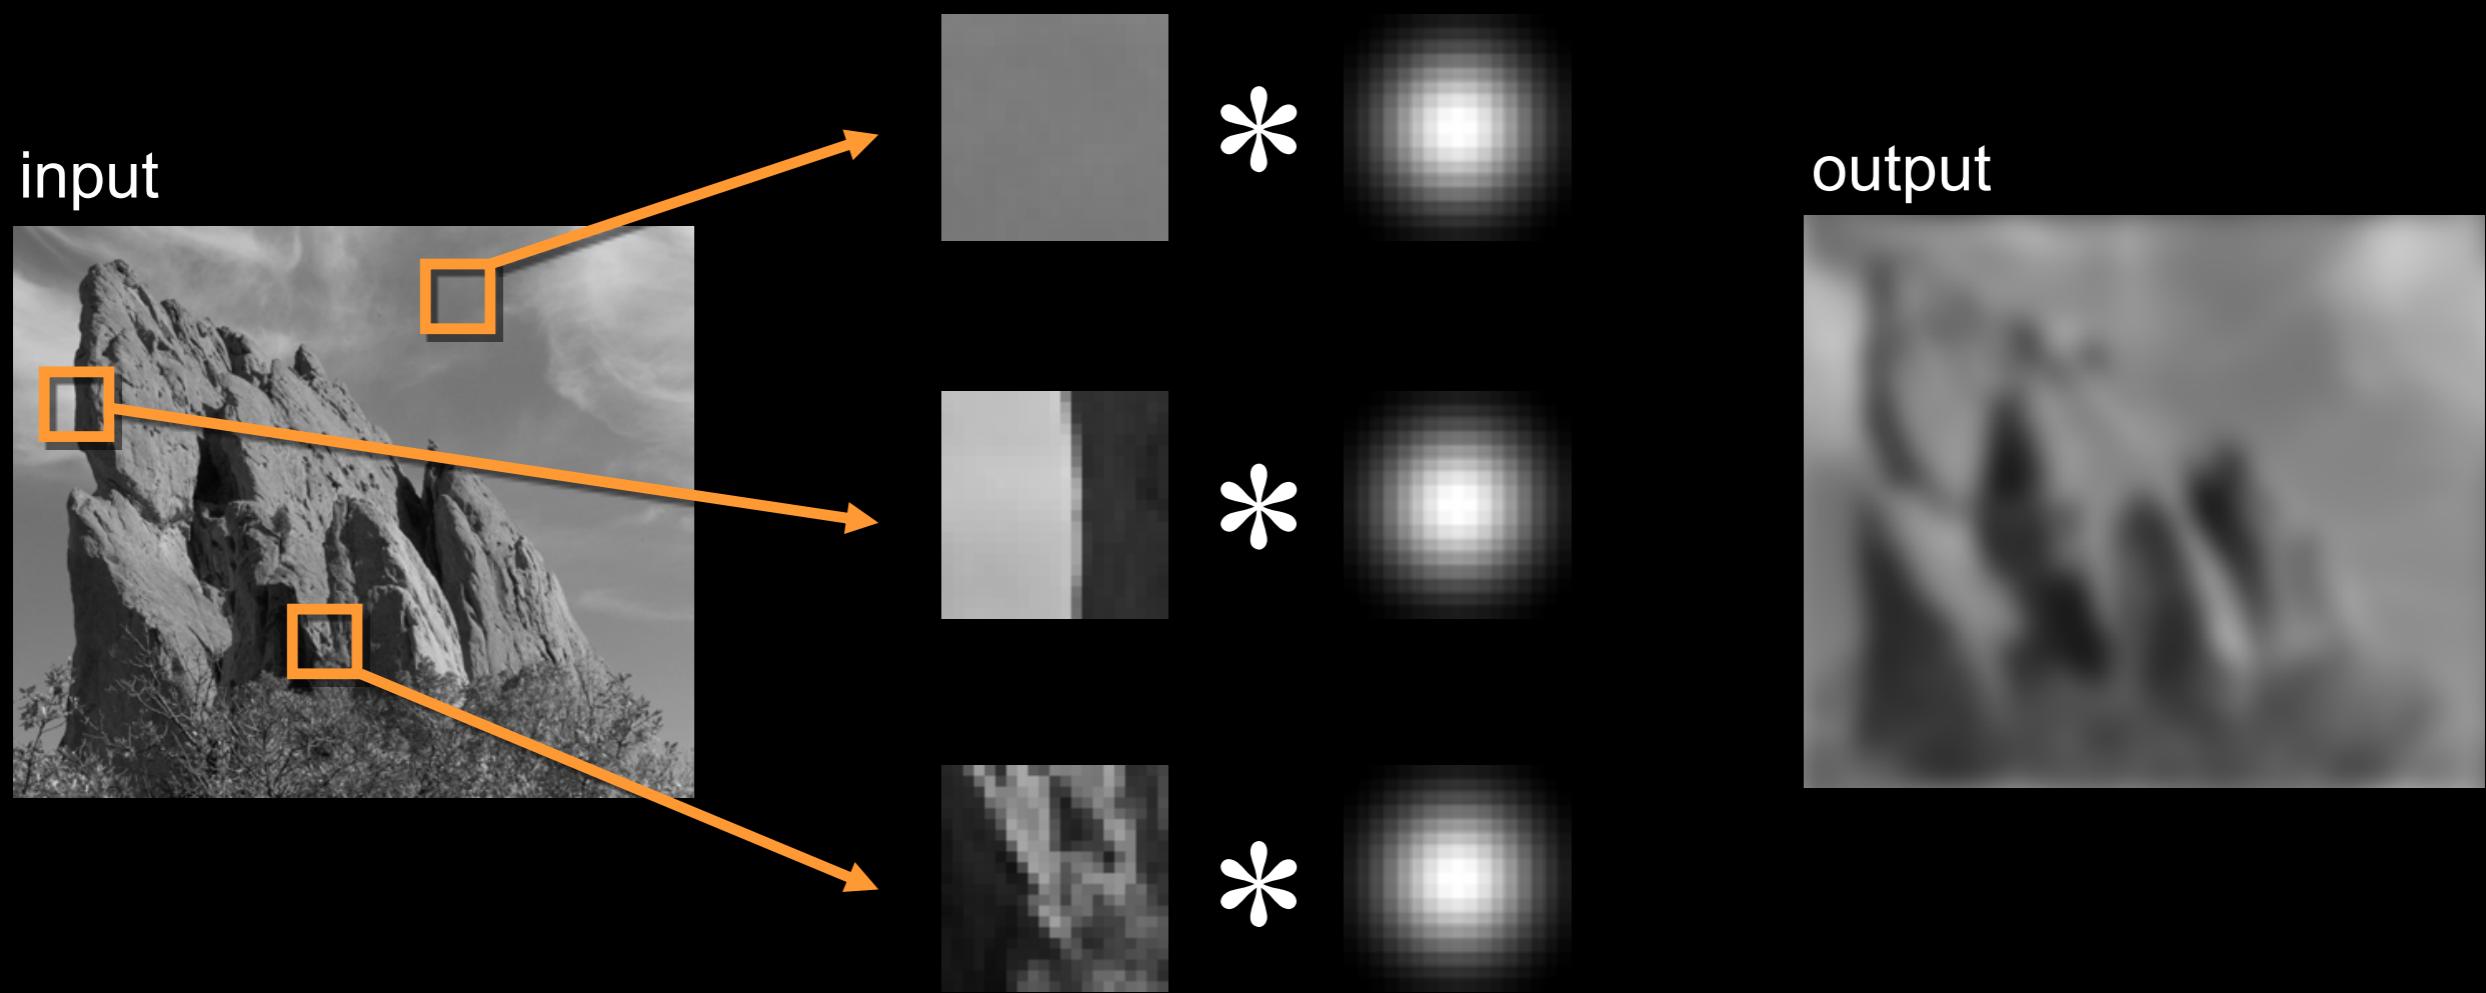
Agenda

- Bilateral filtering
 - The concept
 - Complexity and implementation details
 - Fast bilateral filter – code
 - Bilateral solver
- Non-local means
 - The concept
 - Complexity and implementation details
 - PCA-NLM – code

Bilateral Filter

- [IMPORTANT] Not to be confused with “bilinear” filter that we learned in the previous classes
- Bilinear interpolation – filling the missing values through linear interpolation (which is not in our best interests today).
- Bilateral filter – a non-linear, edge-preserving, noise-reducing filter.

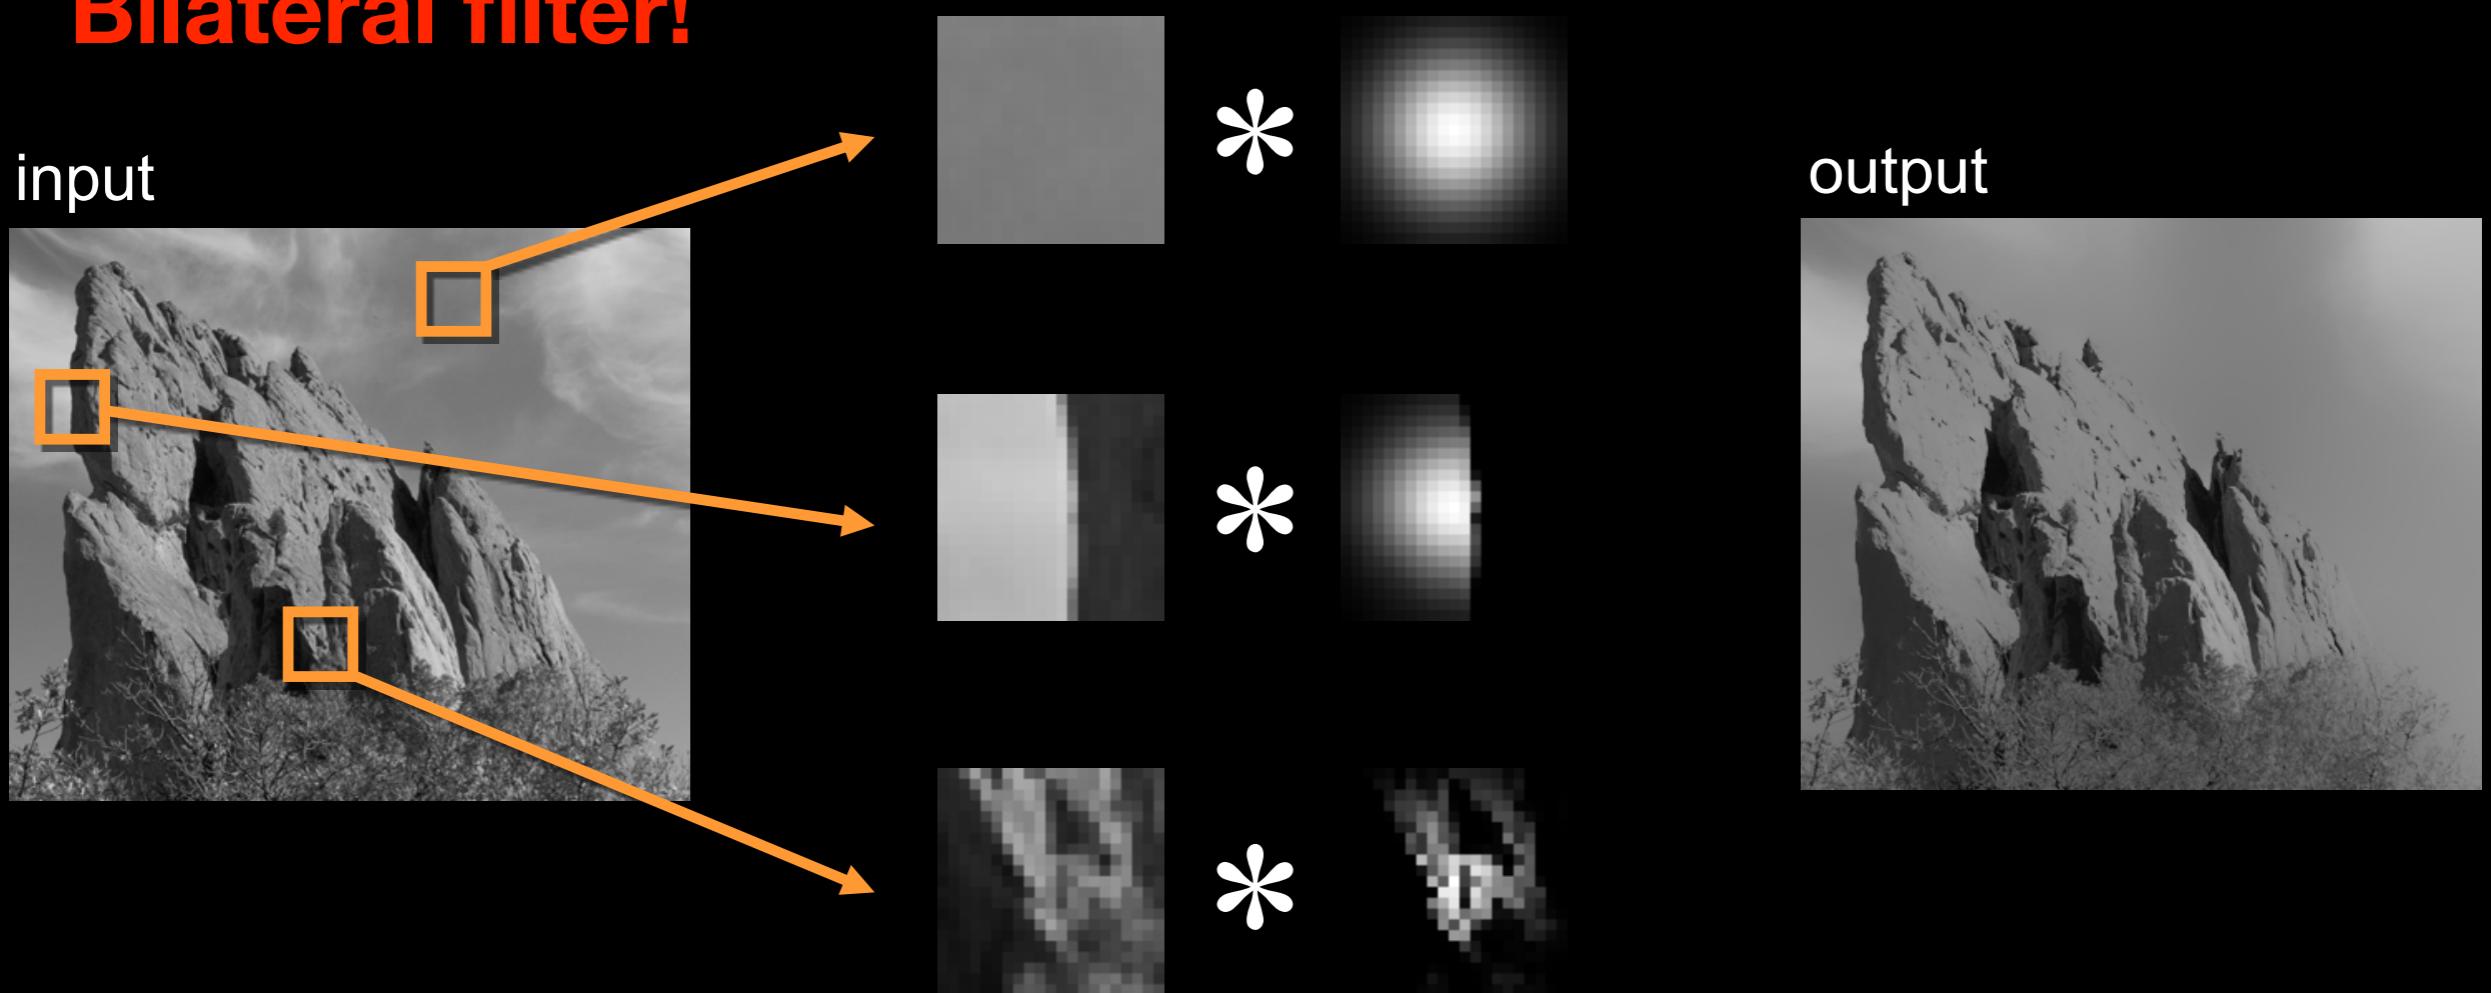
Problem with Gaussian blur?



Same Gaussian kernel everywhere.

What do we rather want?

Bilateral filter!



The kernel shape depends on the image content.

Bilateral filter: definition

Same idea: weighted average of pixels

$$BF[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q) I_q$$

Normalisation

“space” weight “range” weight

$$G_\sigma(x) = e^{-\frac{x^2}{2\sigma^2}}$$

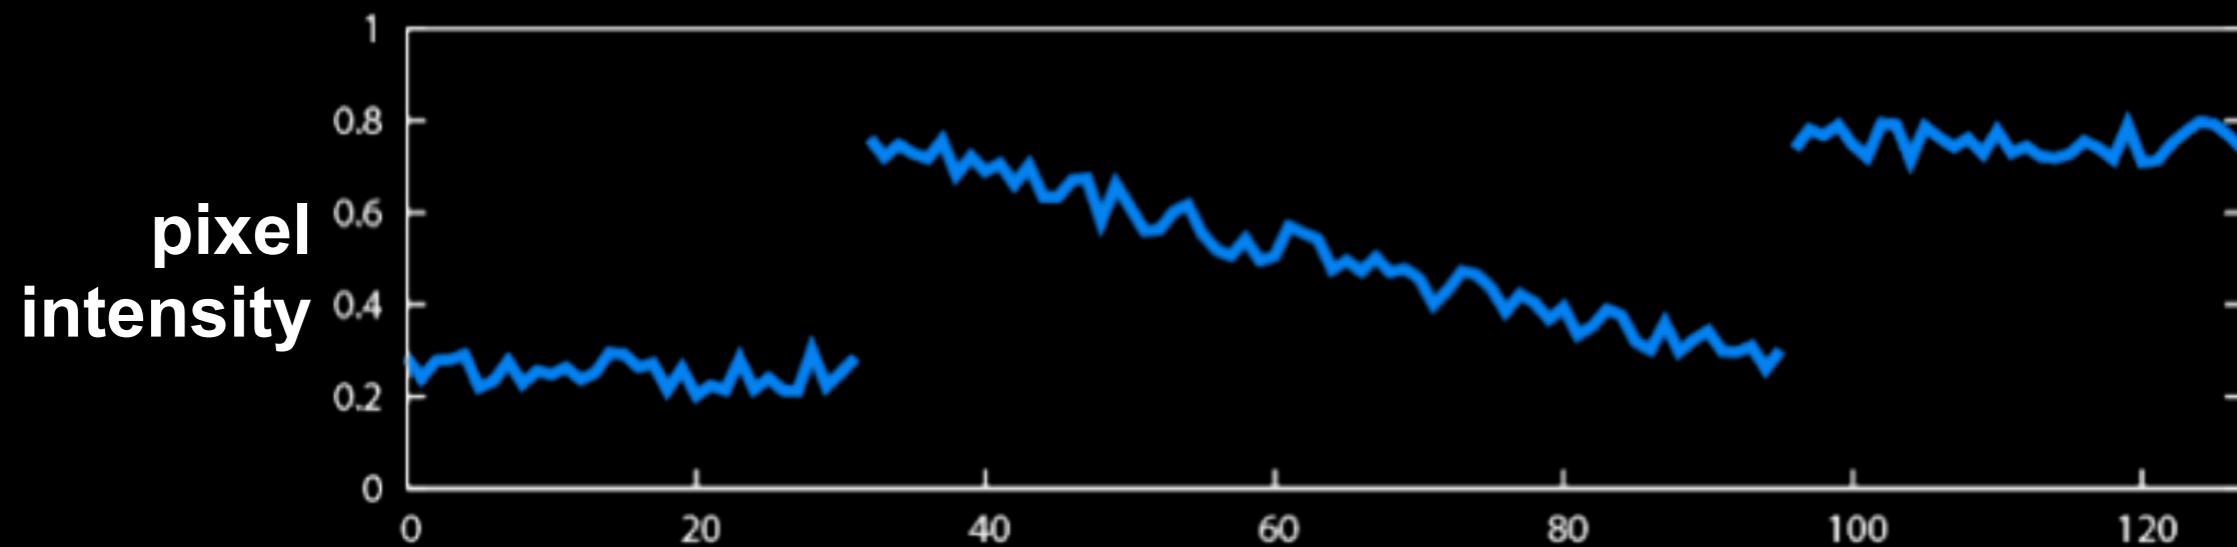
$$W = \sum_{q \in S} G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q)$$

Illustration: 1D

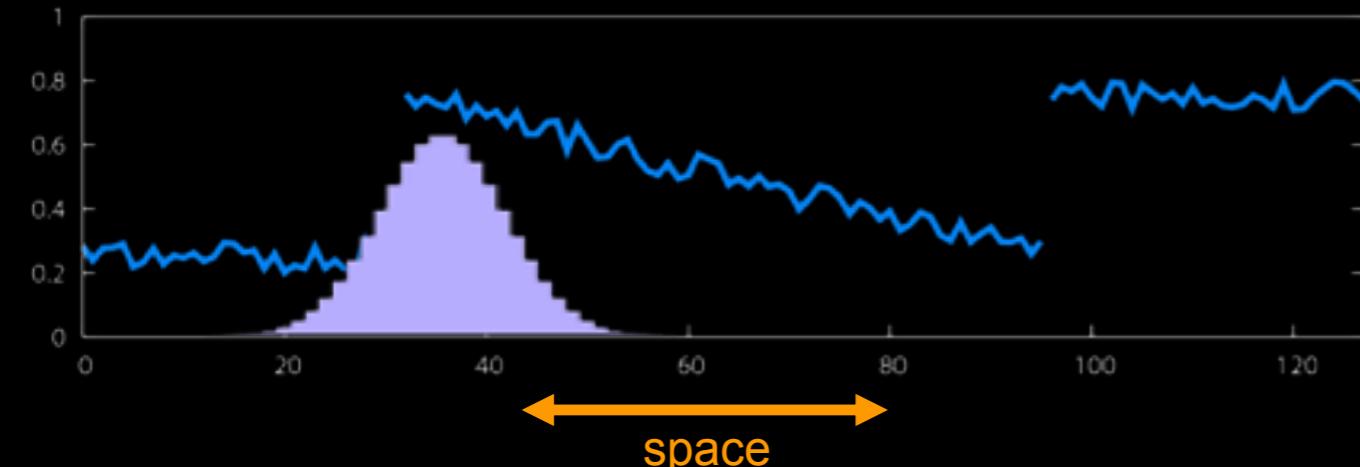
1D image = line of pixels



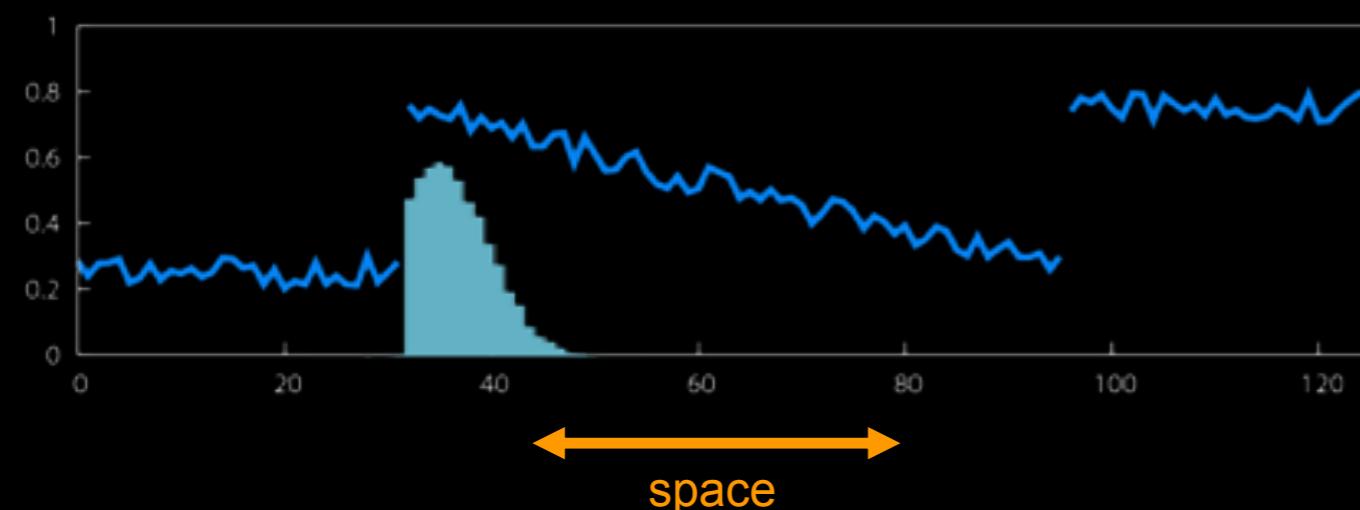
Better visualised as a plot



Gaussian blur vs. Bilateral filter



Gaussian blur



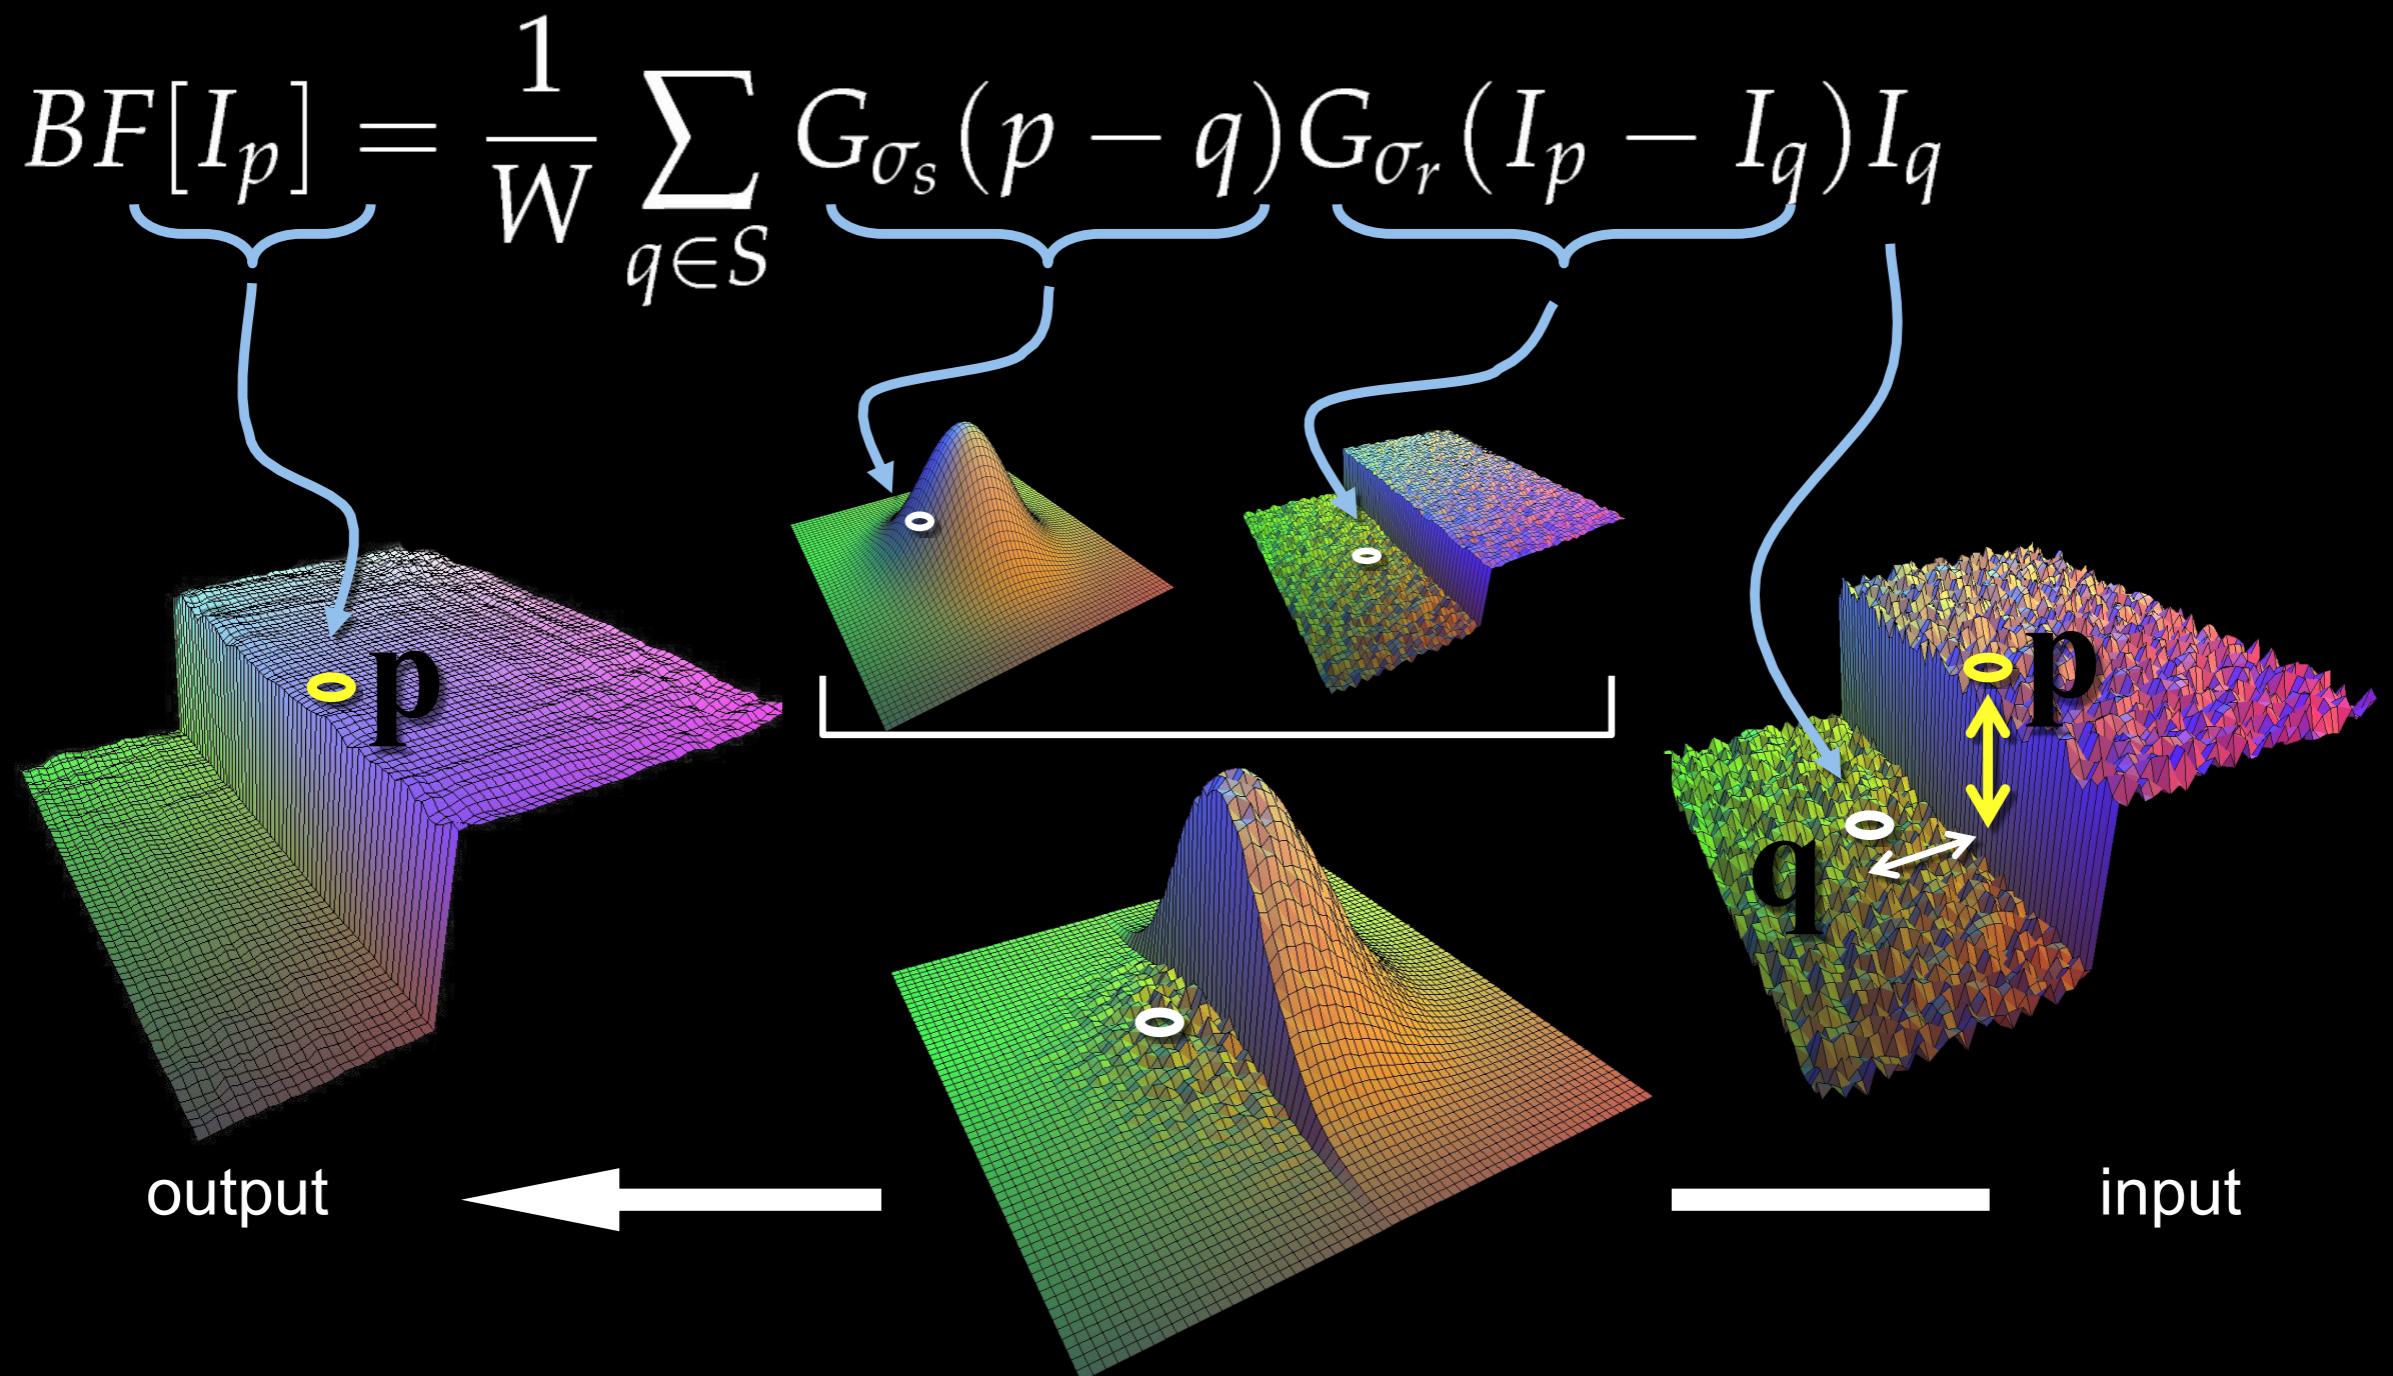
Bilateral filter

$$GB[I_p] = \sum_{q \in S} G_\sigma(p - q) I_q$$

space range

$$BF[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q) I_q$$

Bilateral filter on a height field



Space and range hyper-parameters

$$BF[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q) I_q$$

space σ_s : spatial extent of the kernel, size of the considered neighbourhood.

range σ_r : the estimated standard deviation of the noise.

Complexity

$$BF[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q) I_p$$

- for “each” pixel p
 - for “each” pixel q
 - compute: $G_{\sigma_s}(p - q) G_{\sigma_r}(I_p - I_q)$

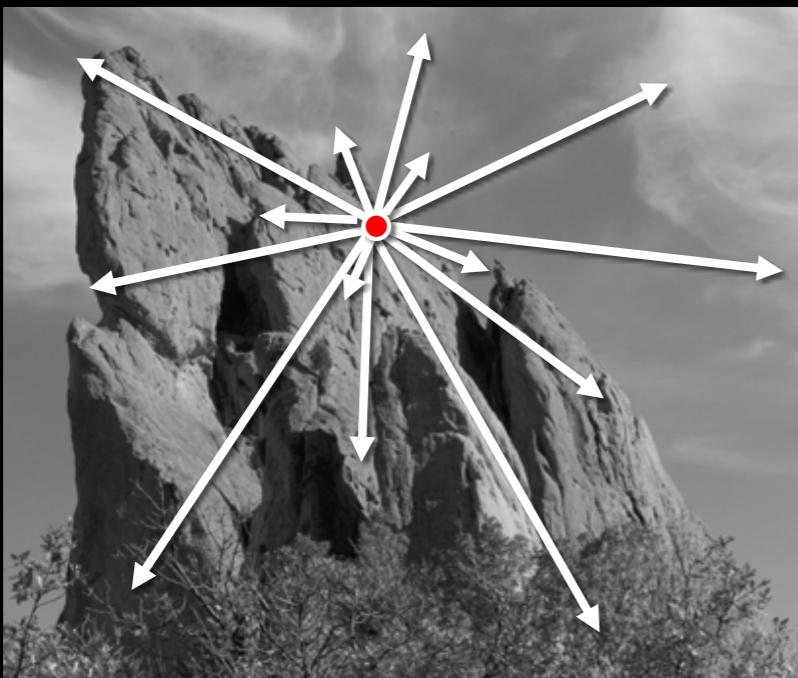
VERY SLOW!!

**For an 8 MP image - takes
64,000,000,000,000 iterations**

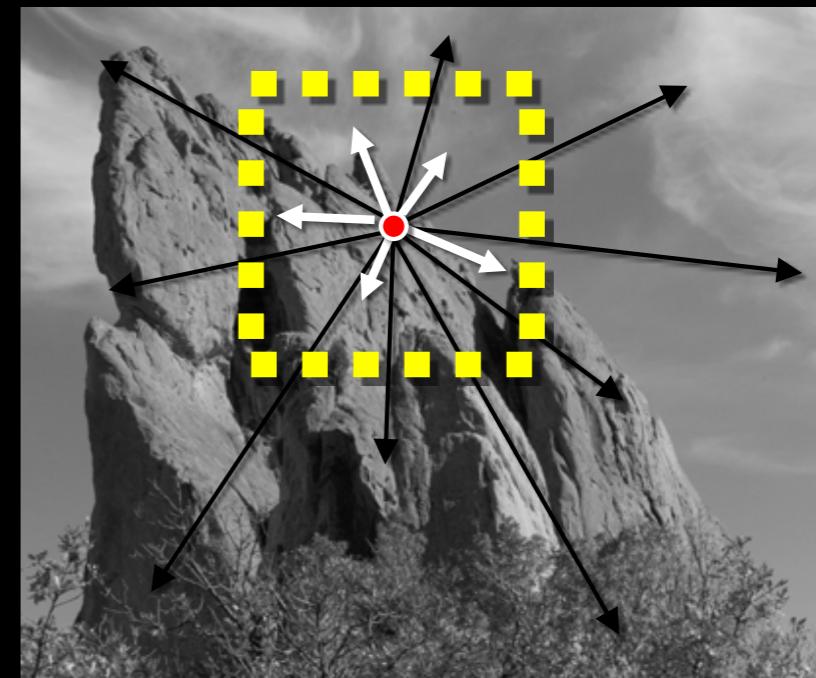
Better bilateral filter

Idea: Far away pixels are negligible, so don't see them at all

looking at all pixels

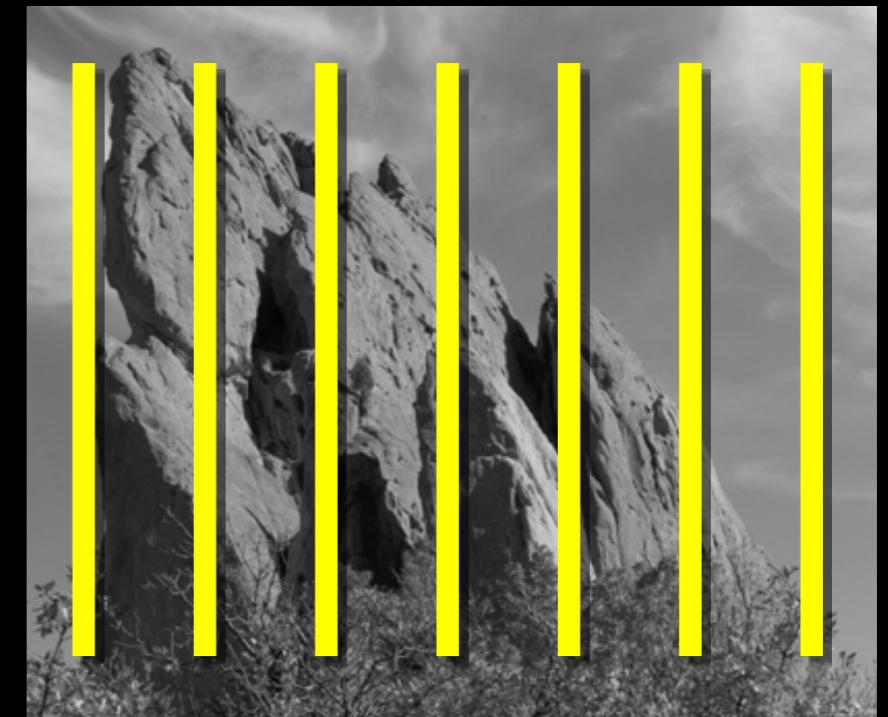
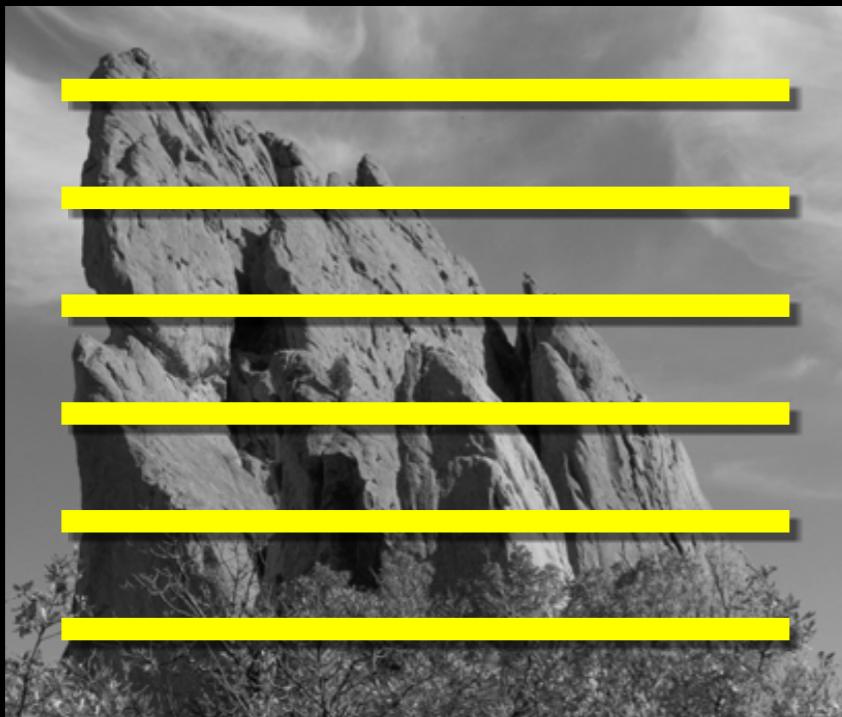


looking at neighbors only



Separable filter

Idea: filter the rows then the columns



Two “cheap” 1D filters instead of an “expensive” 2D filter

Applications: denoising

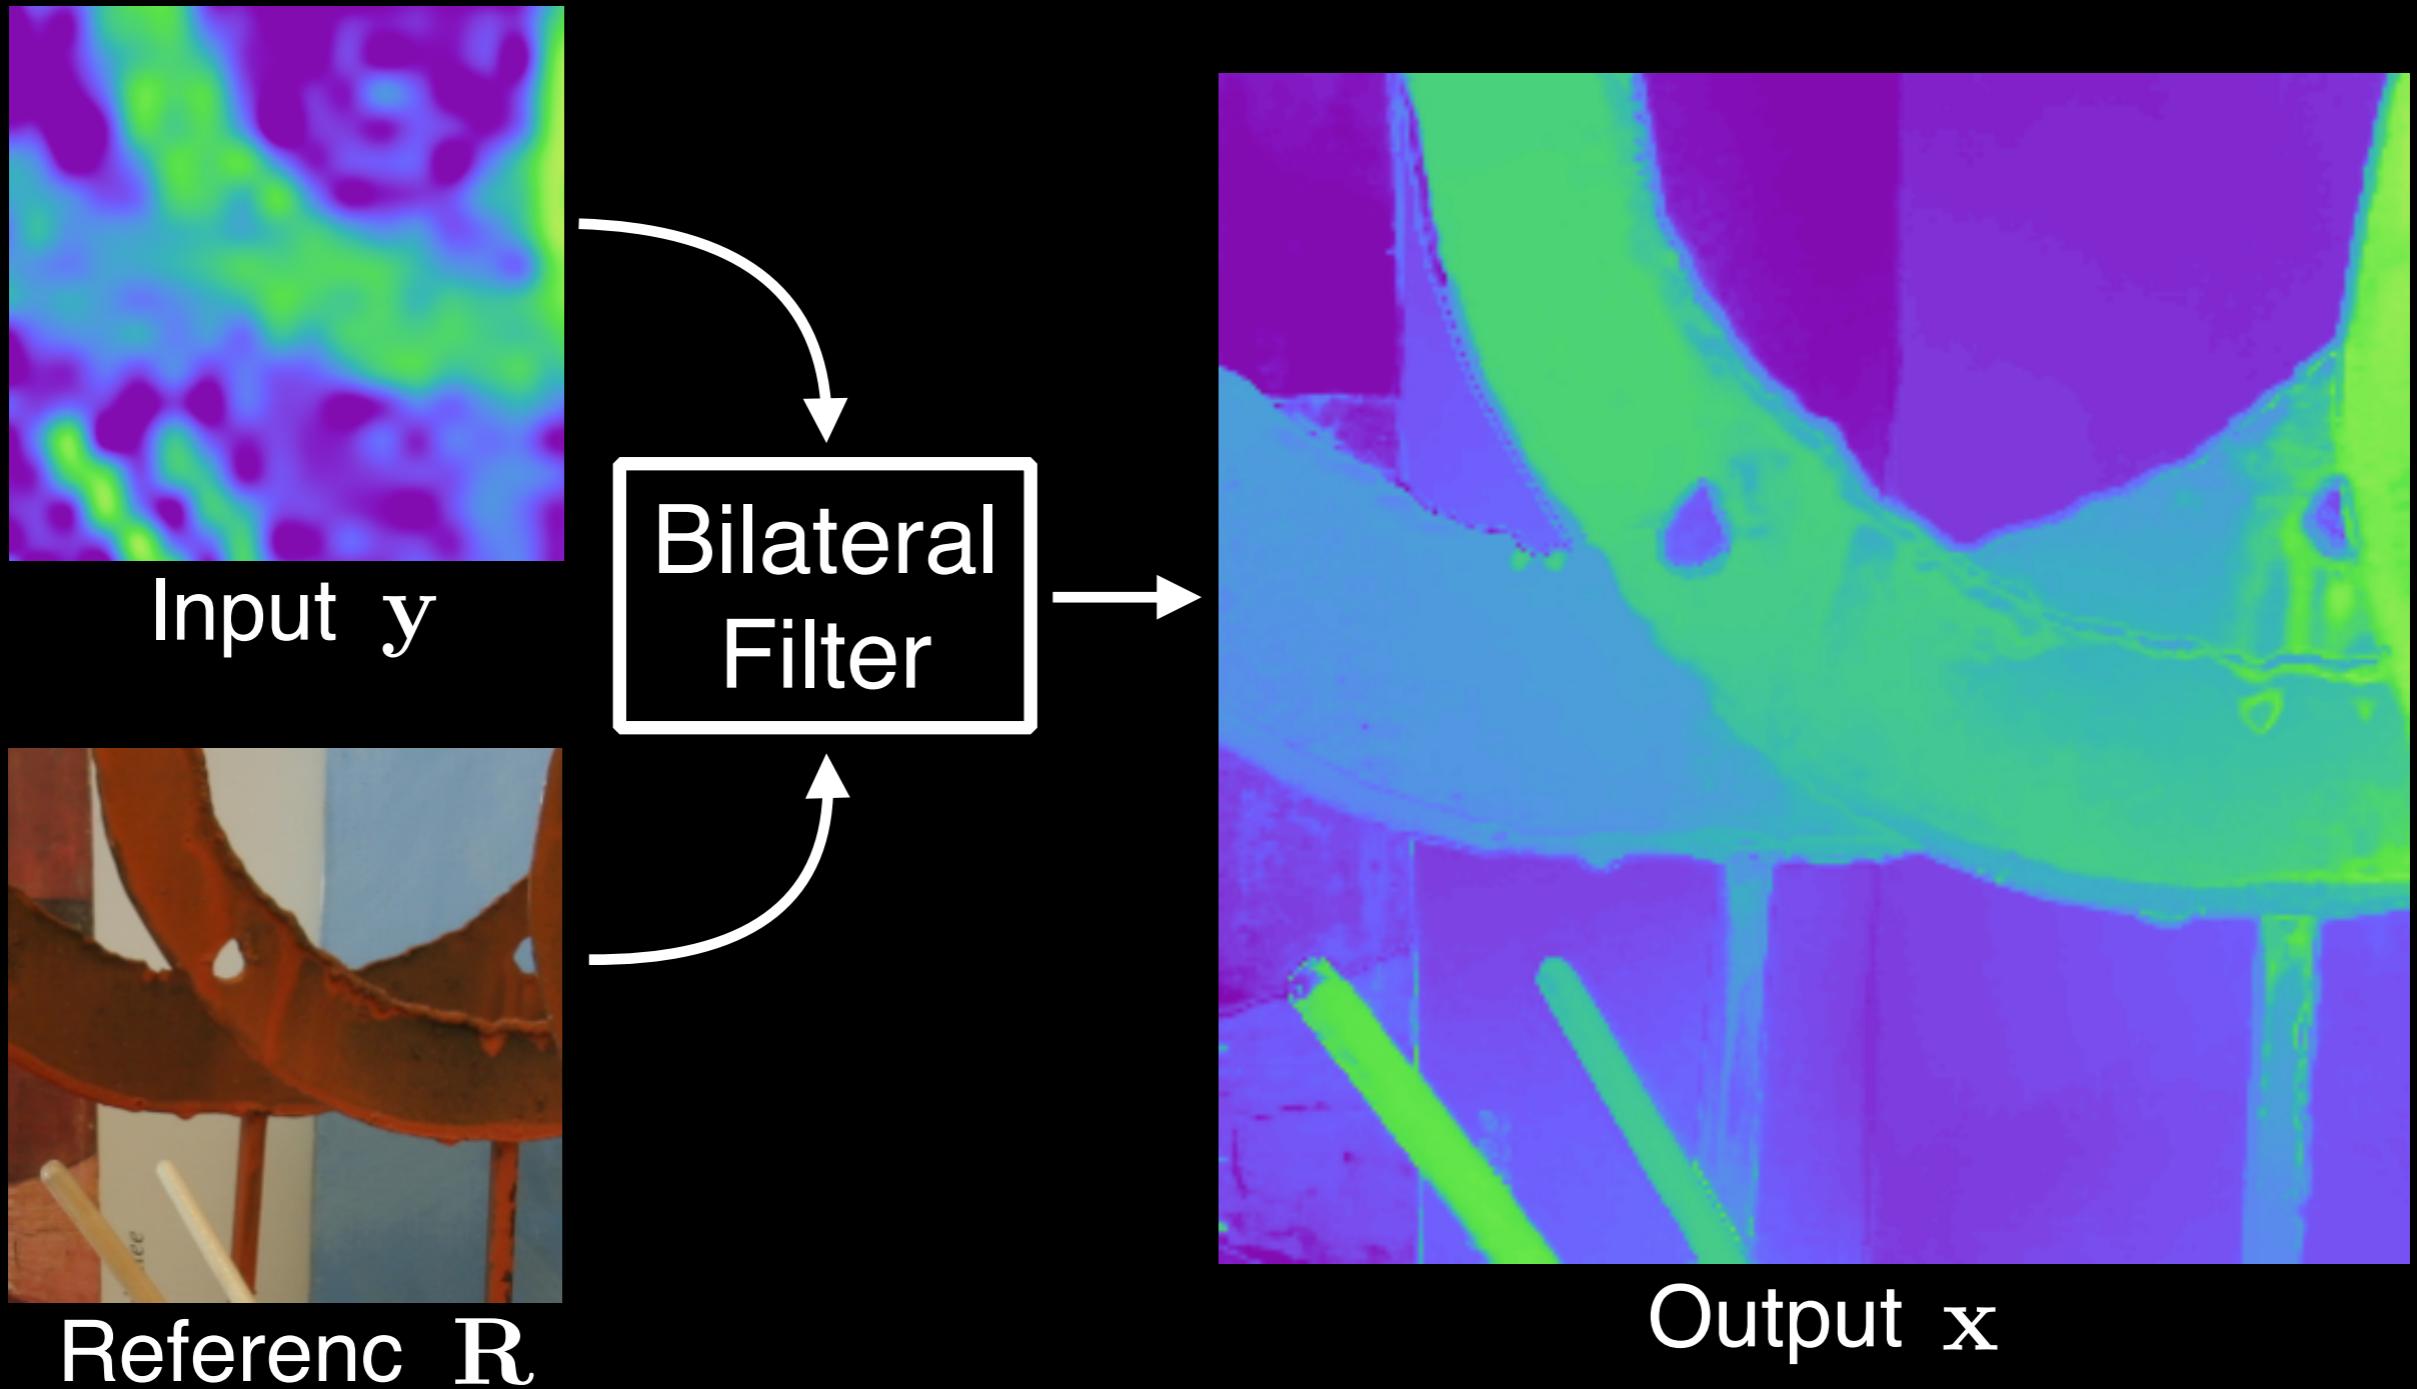
Noisy input



Bilateral filter 7x7 window

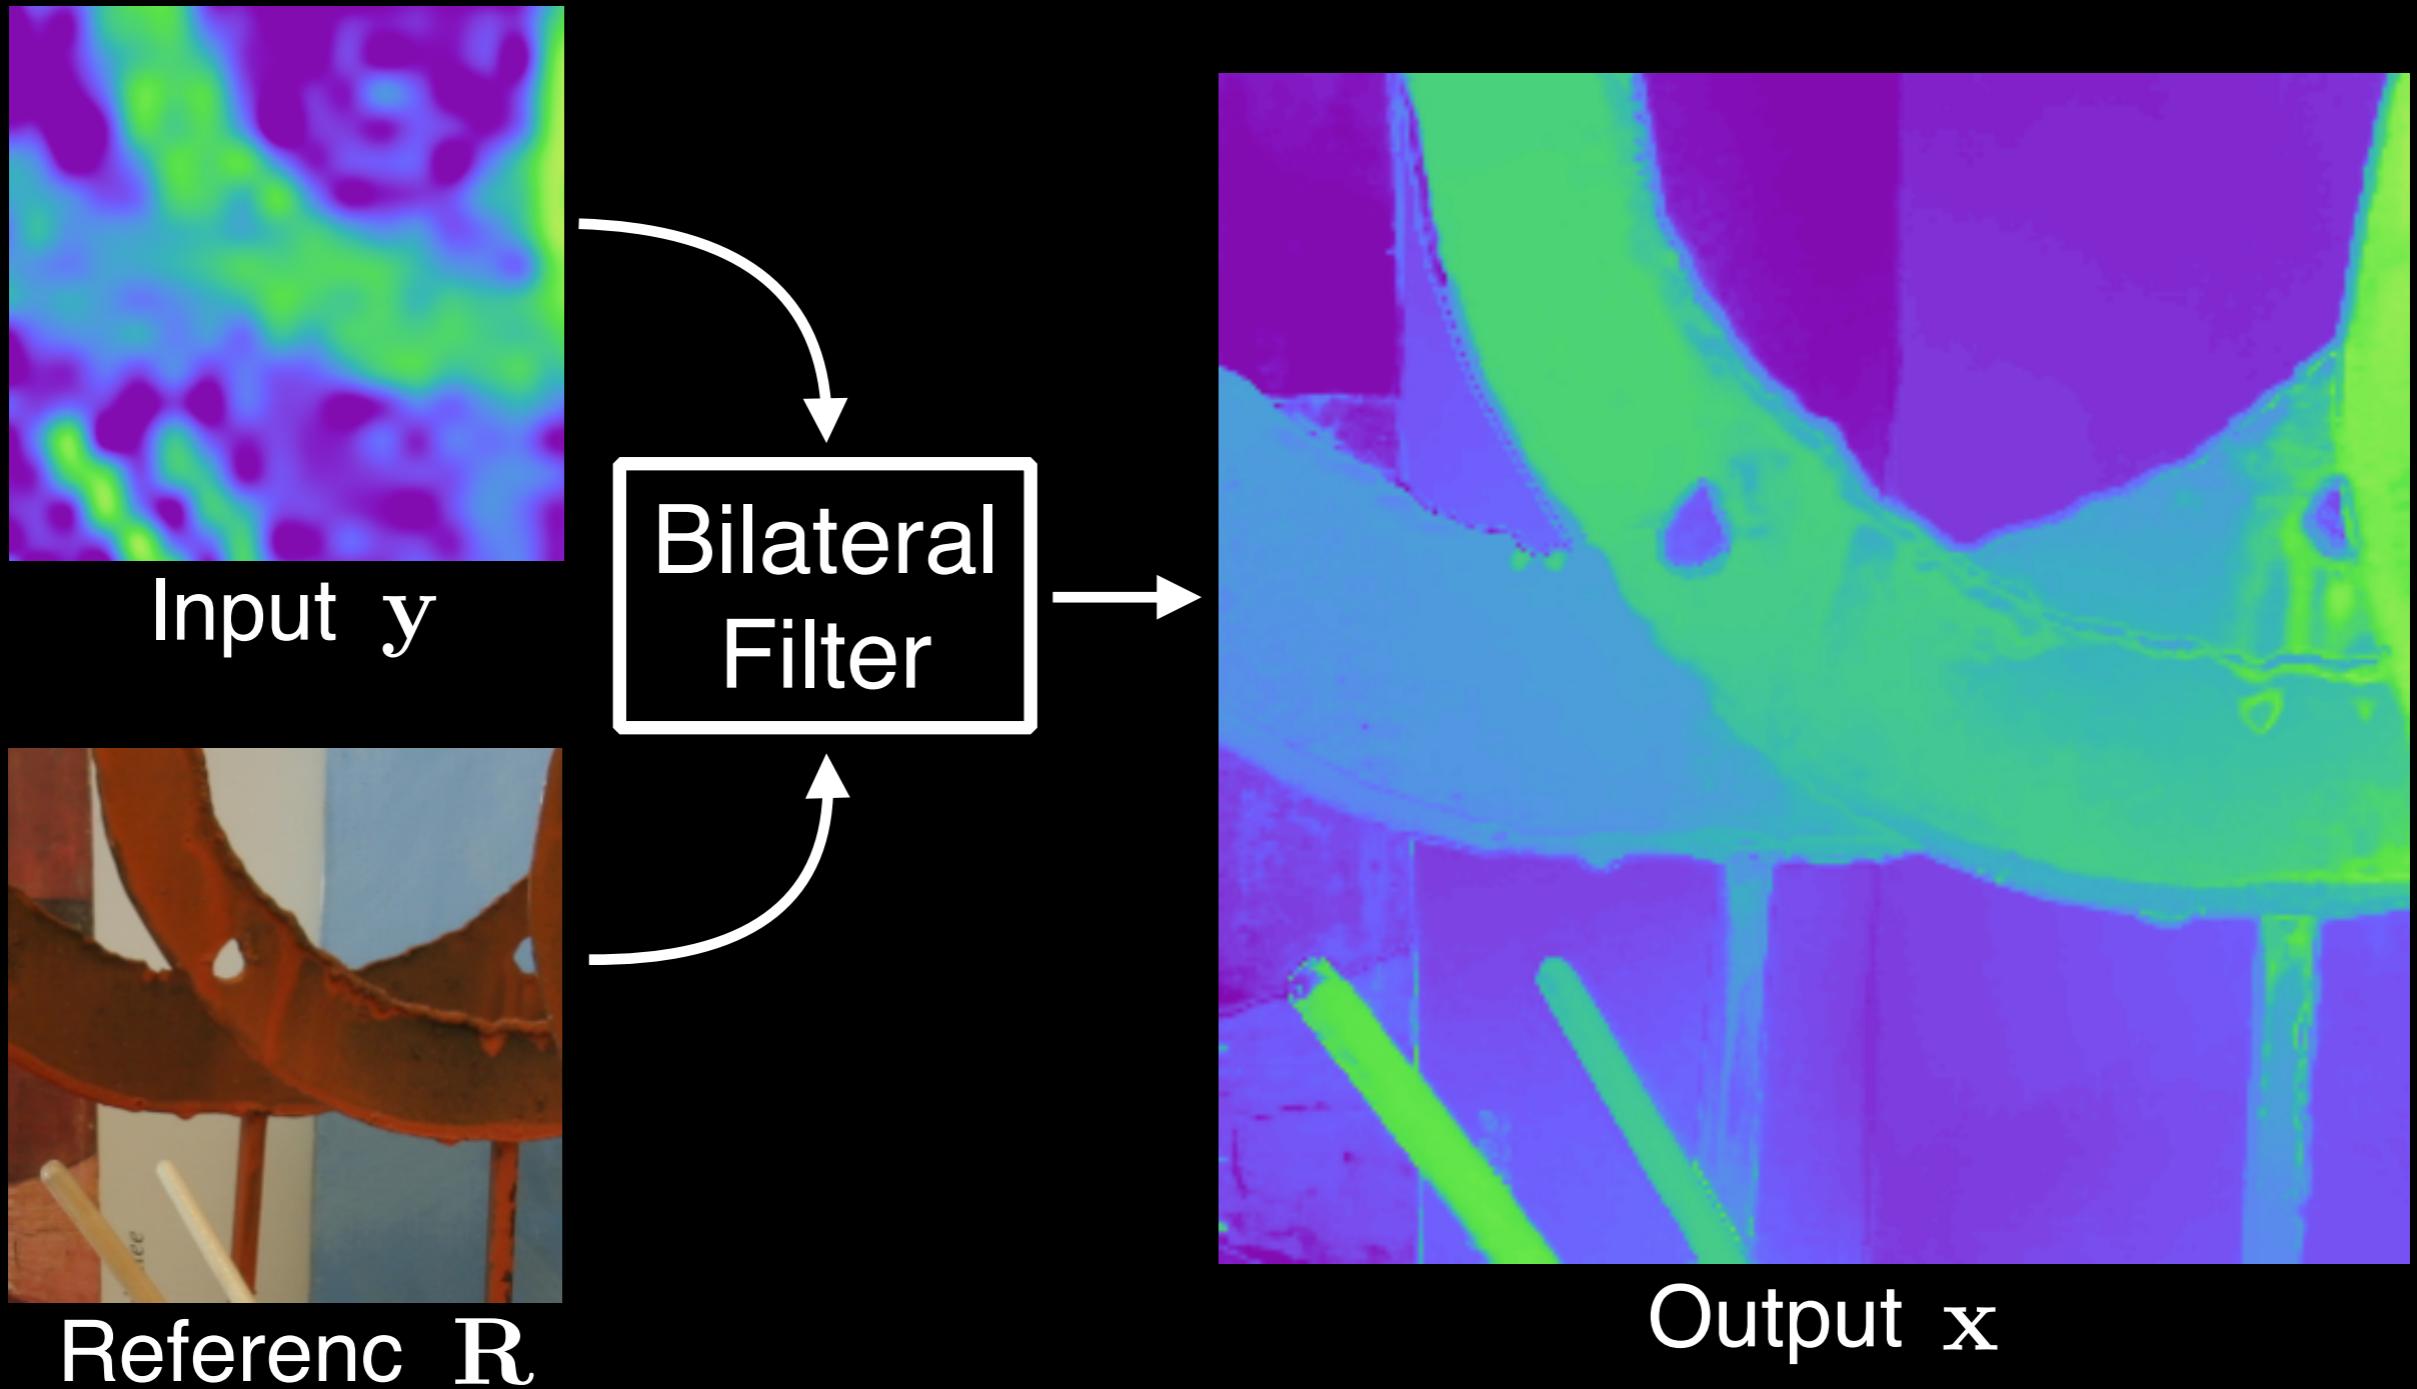


Semantic segmentation



Blur the input with respect to the edges in the reference image.

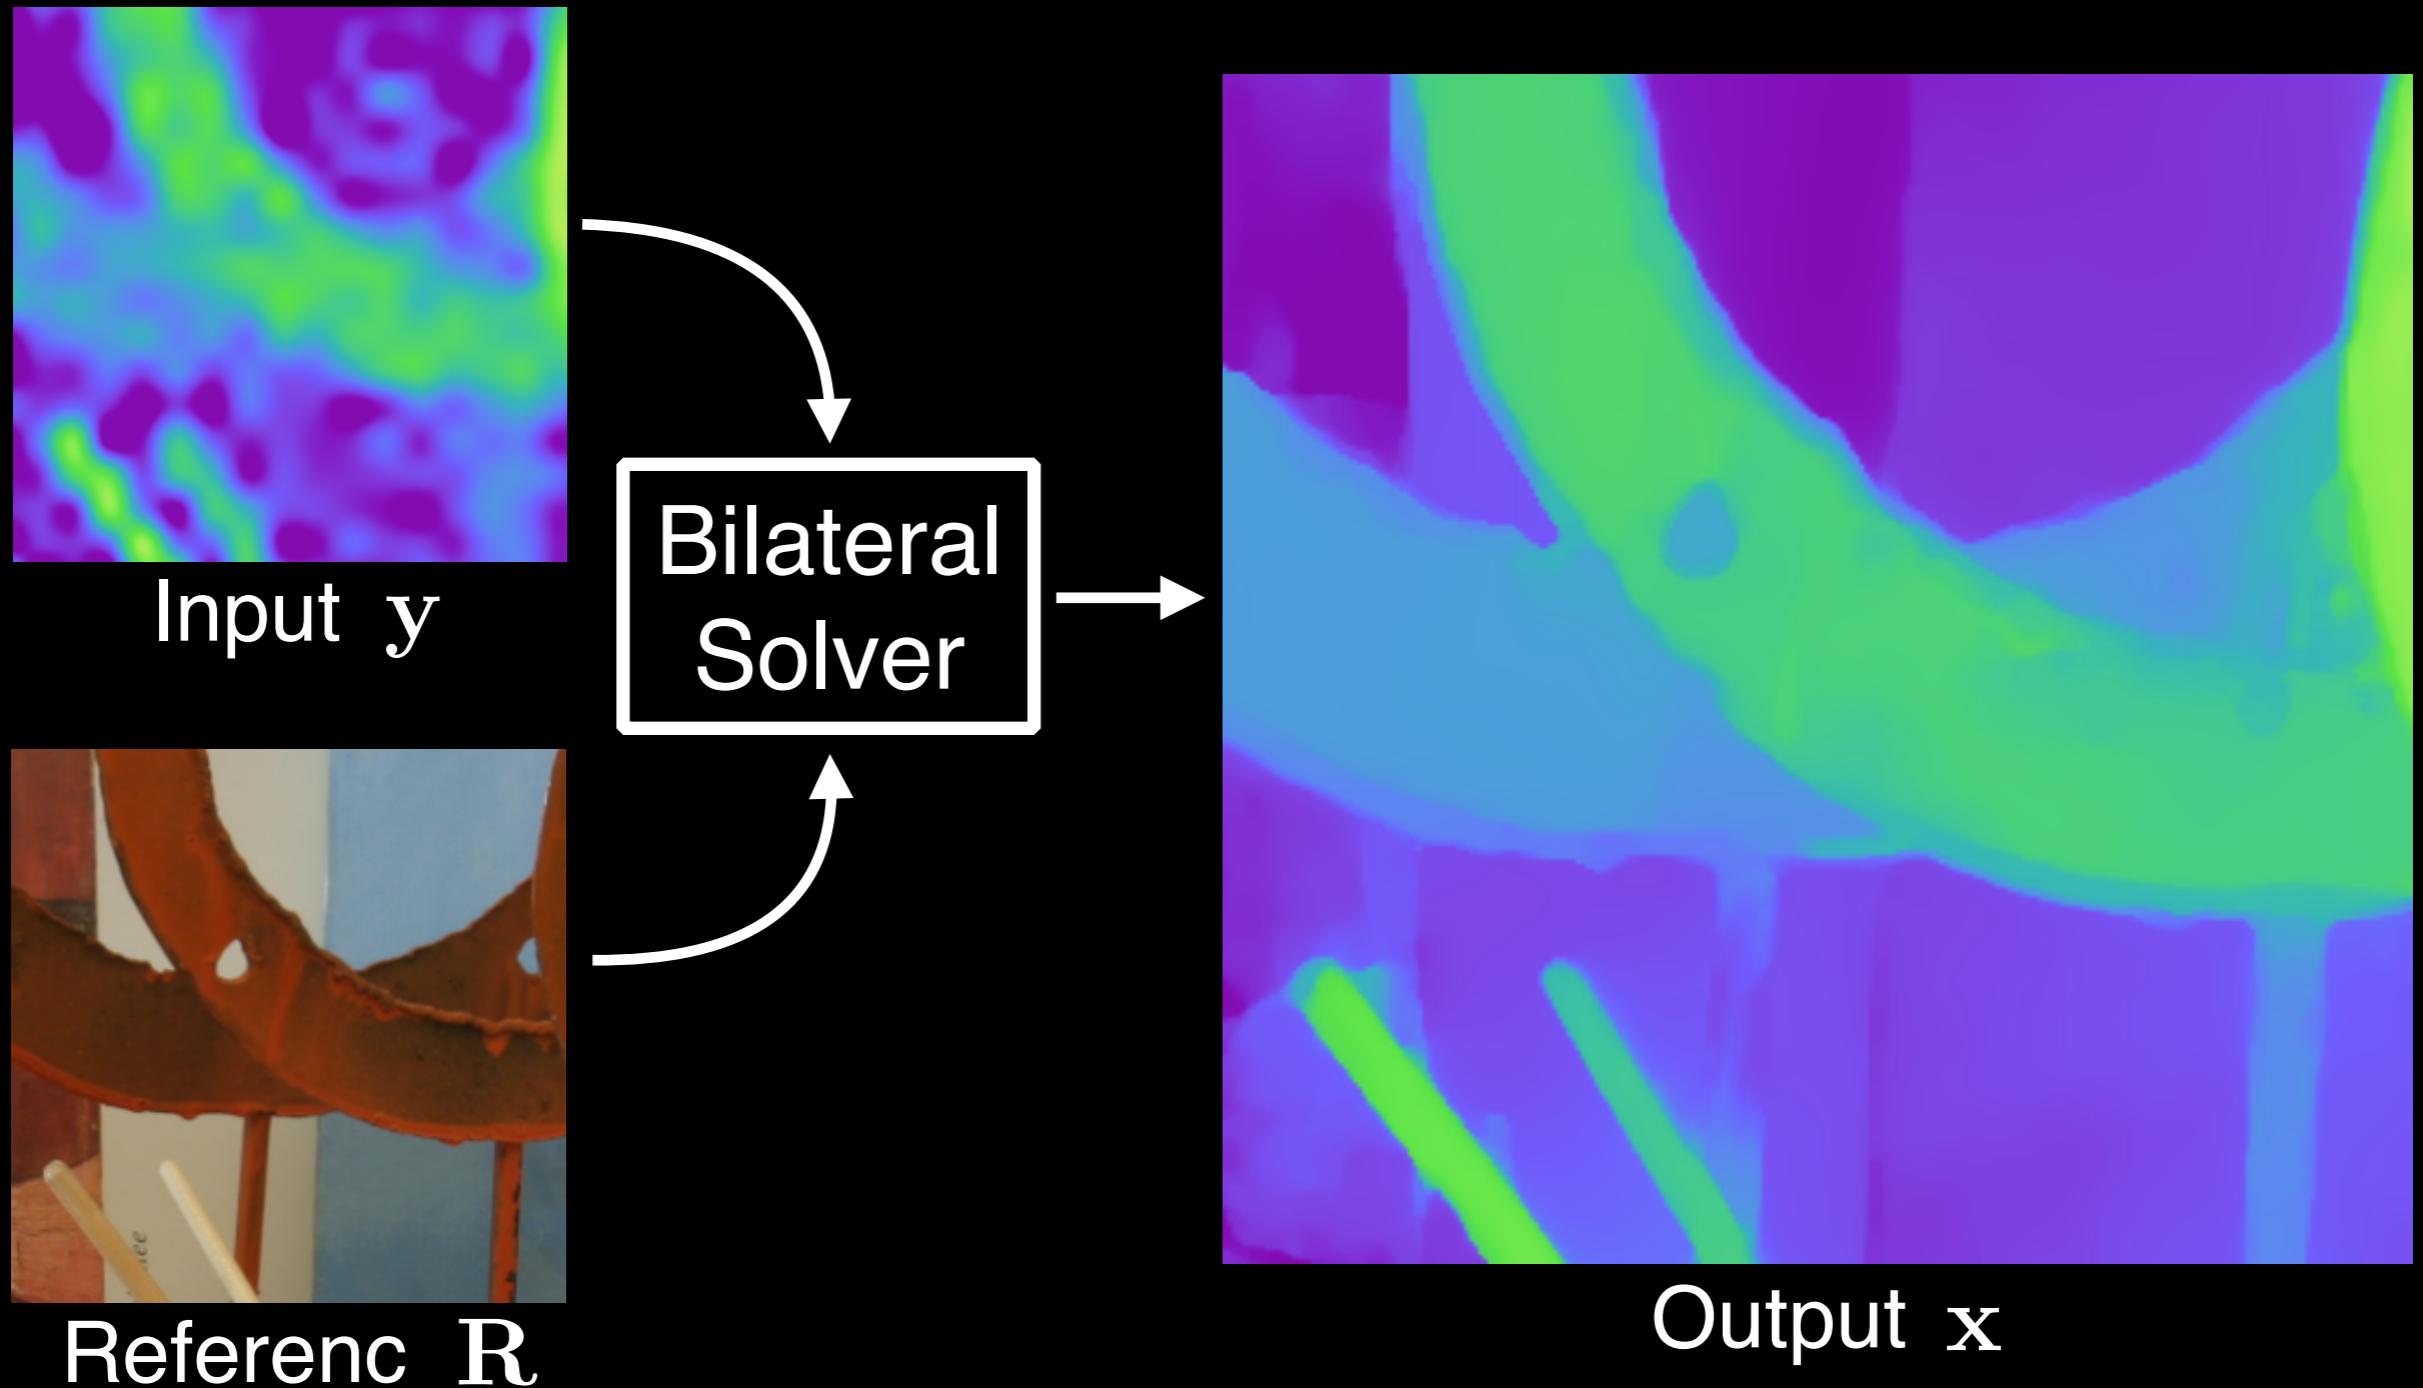
Semantic segmentation



$$\mathbf{x} \leftarrow \mathbf{W}\mathbf{y}$$

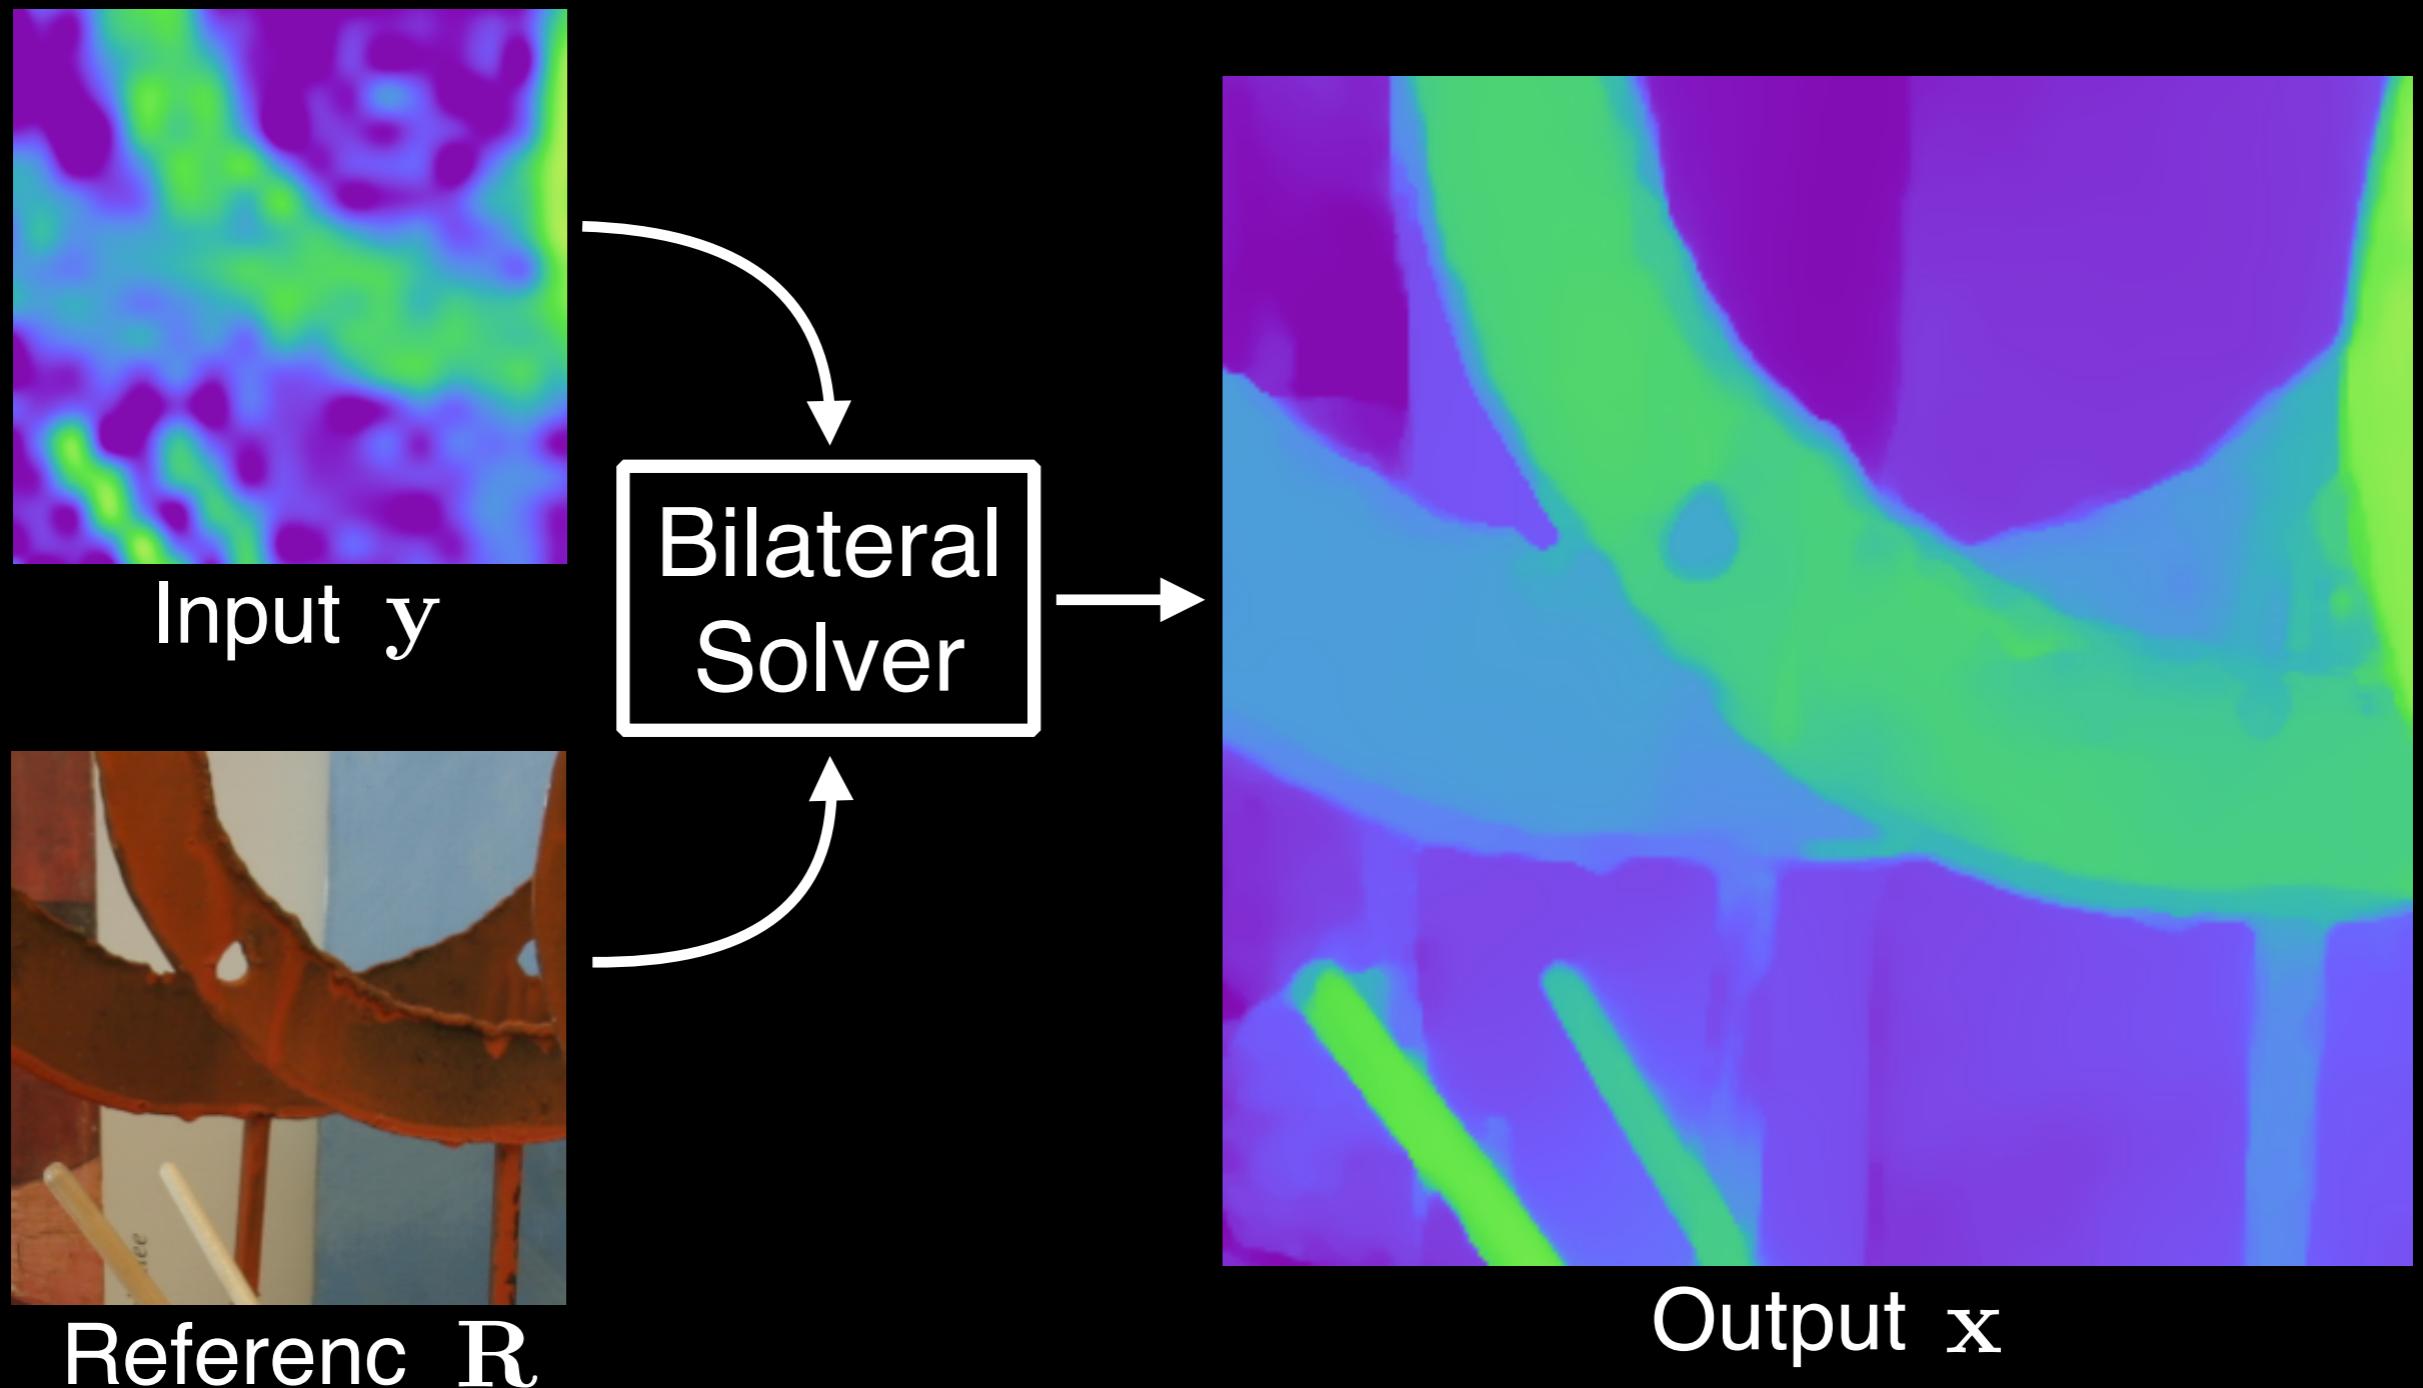
$$W_{i,j} = \exp\left(-\sum_d \frac{(R_{i,d} - R_{j,d})^2}{2\sigma_d^2}\right)$$

Bilateral Solver



Find the image that is as smooth as possible with respect to the reference image, and as close as possible to the input.

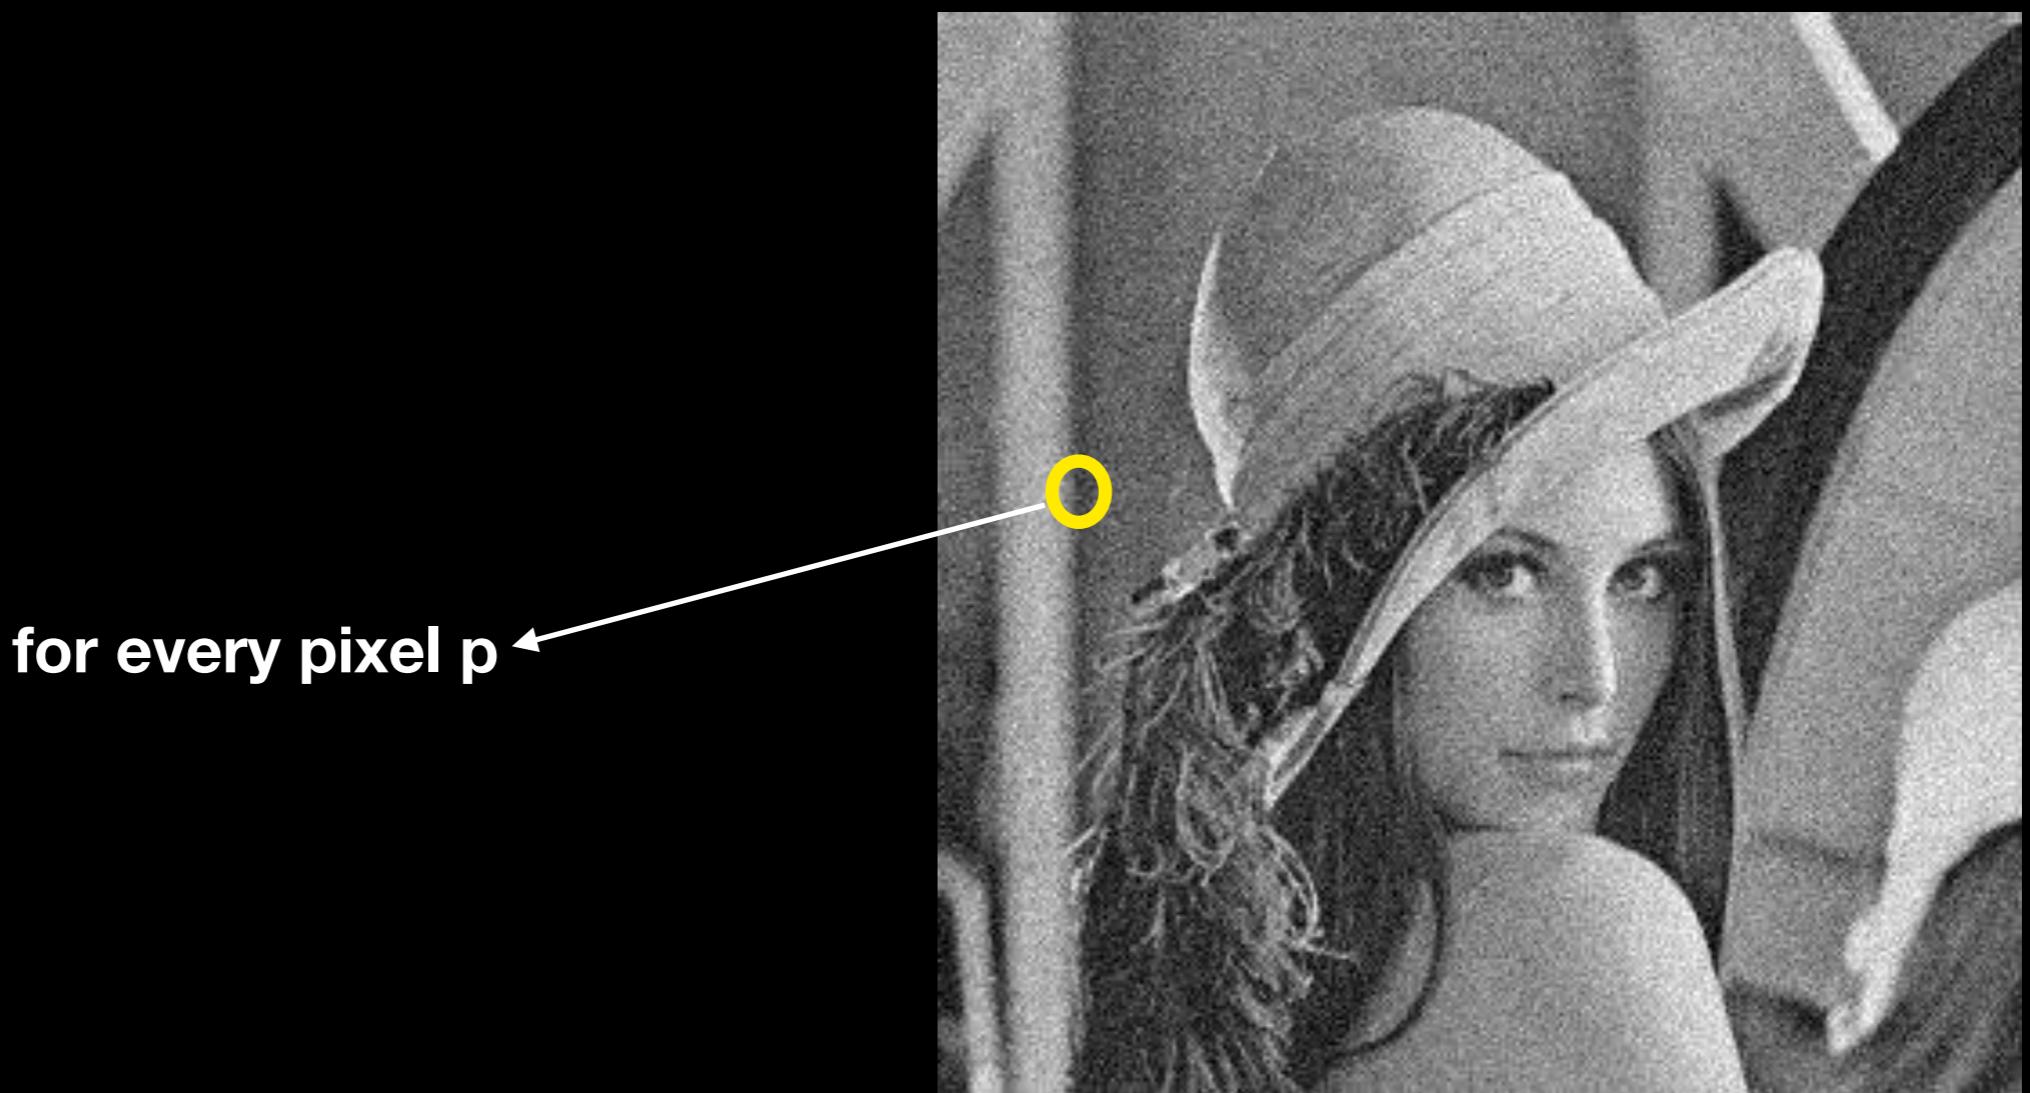
Bilateral Solver



$$\mathbf{x} \leftarrow \arg \min_{\mathbf{x}} \frac{\lambda}{2} \sum_{i,j} W_{i,j} (x_i - x_j)^2 + \sum_i (x_i - y_i)^2$$

Non-local means

Bilateral filter: Average neighbours with similar intensities
Non-local means: Average neighbours with similar neighbourhoods



Non-local means

Bilateral filter: Average neighbours with similar intensities

Non-local means: Average neighbours with similar neighbourhoods!

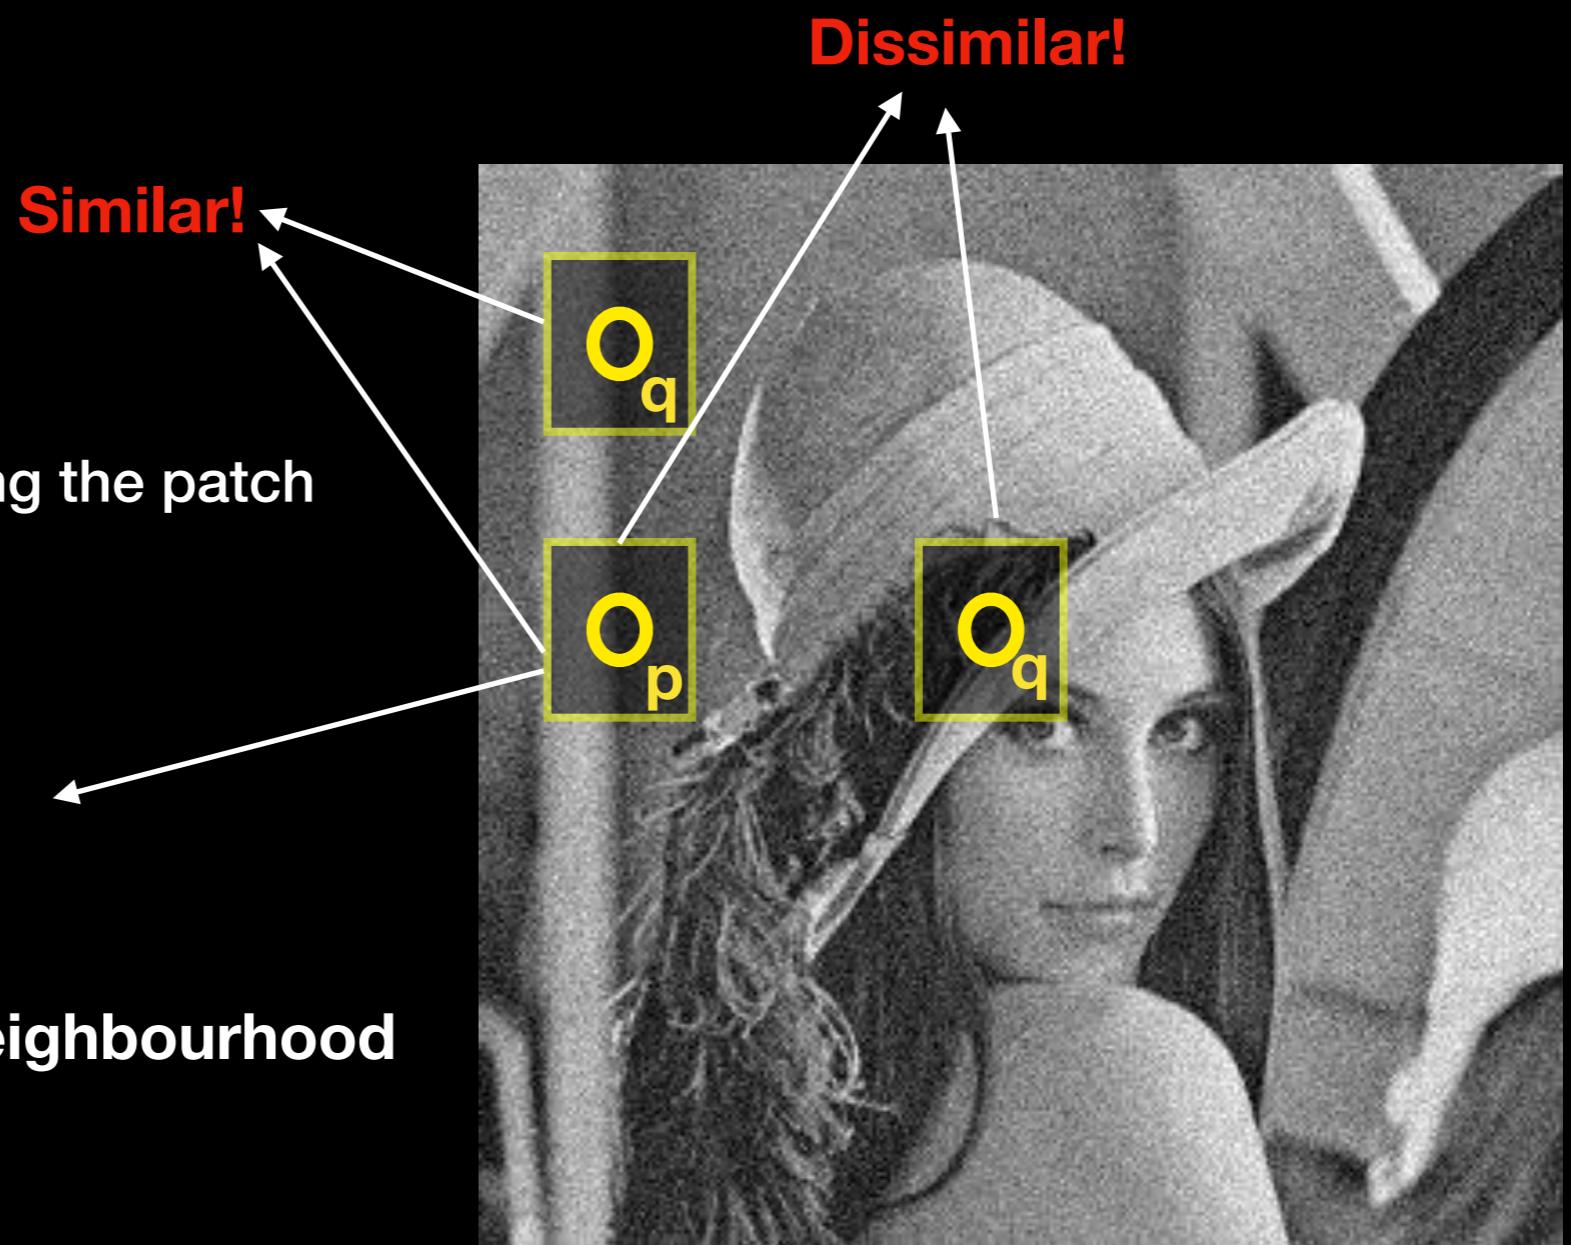
Filter with this!

$$\|v_p - v_q\|^2$$

v_p, v_q are vectors describing the patch

for every pixel p

- define a small neighbourhood



Non-local means

Filter with this!

$$\downarrow \\ ||v_p - v_q||^2$$

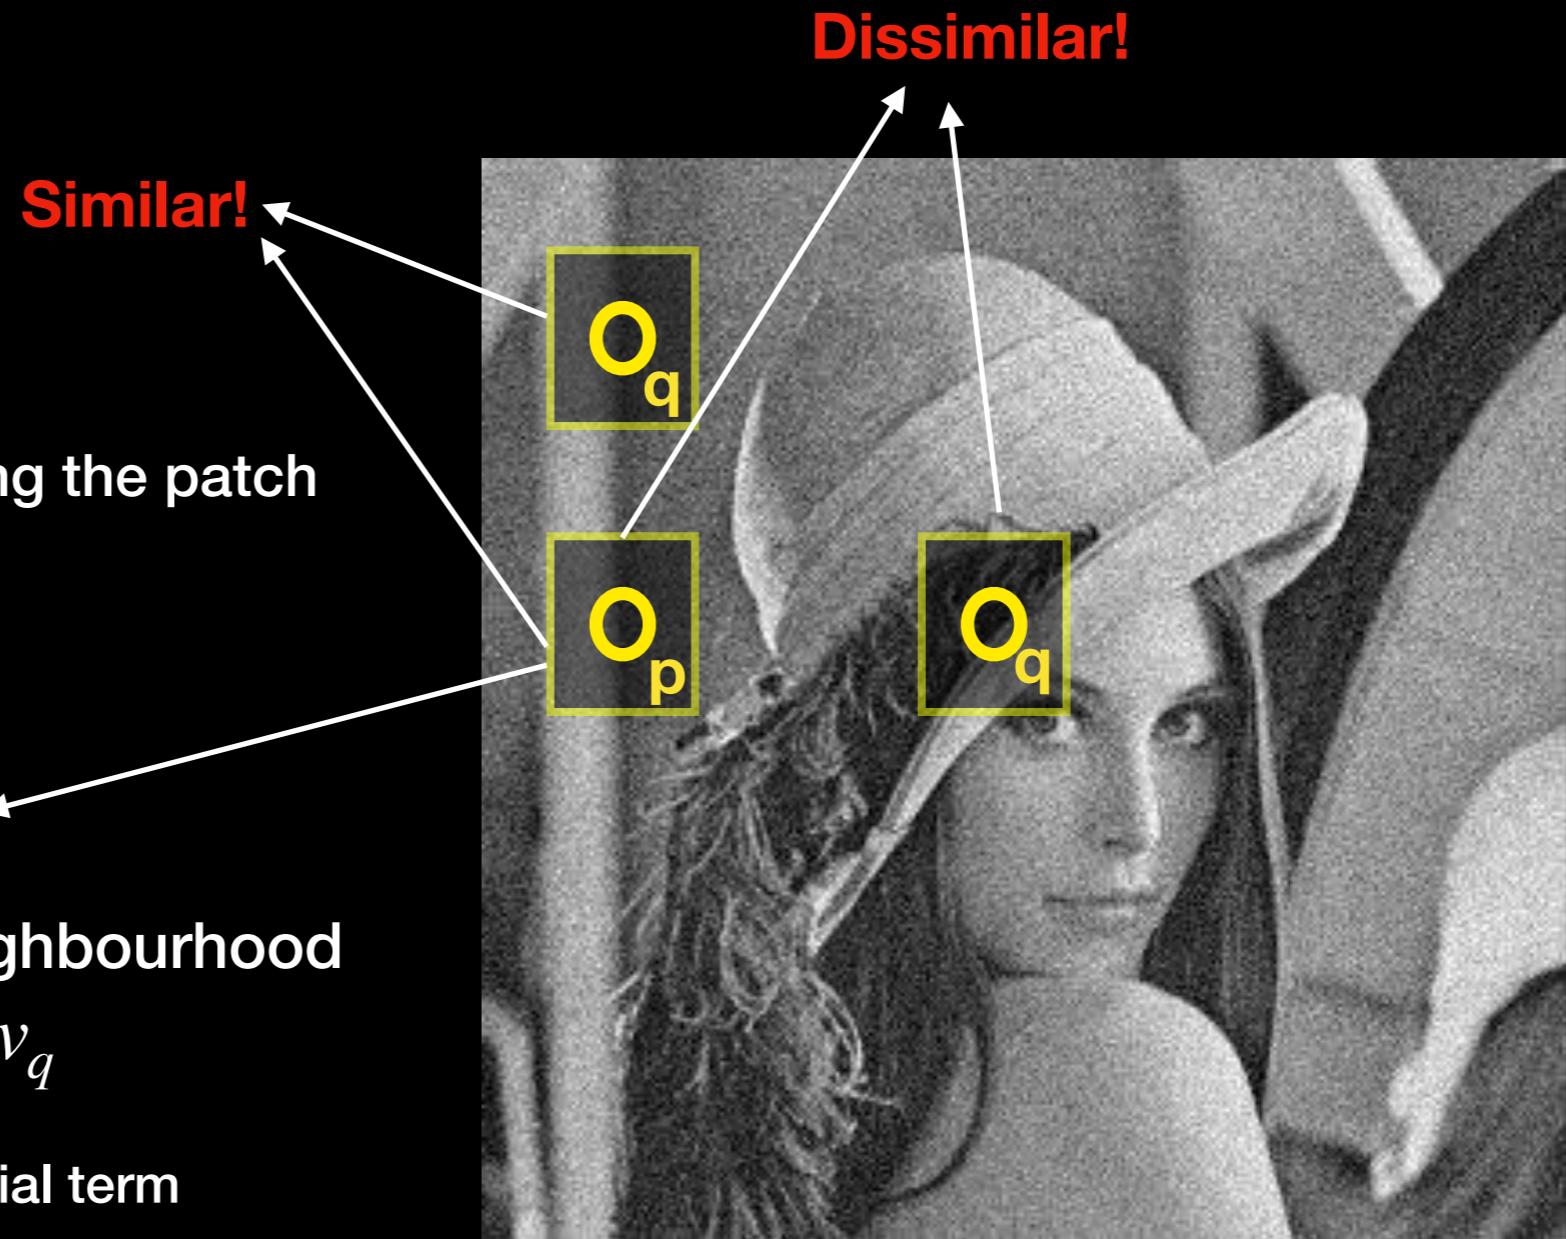
v_p, v_q are vectors describing the patch

for every pixel p

- define a small neighbourhood
- define vectors v_p, v_q

No spatial term

$$NLM[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_s}(p-q) G_{\sigma_r}(||v_p - v_q||^2) I_p$$



Non-local means

Filter with this!

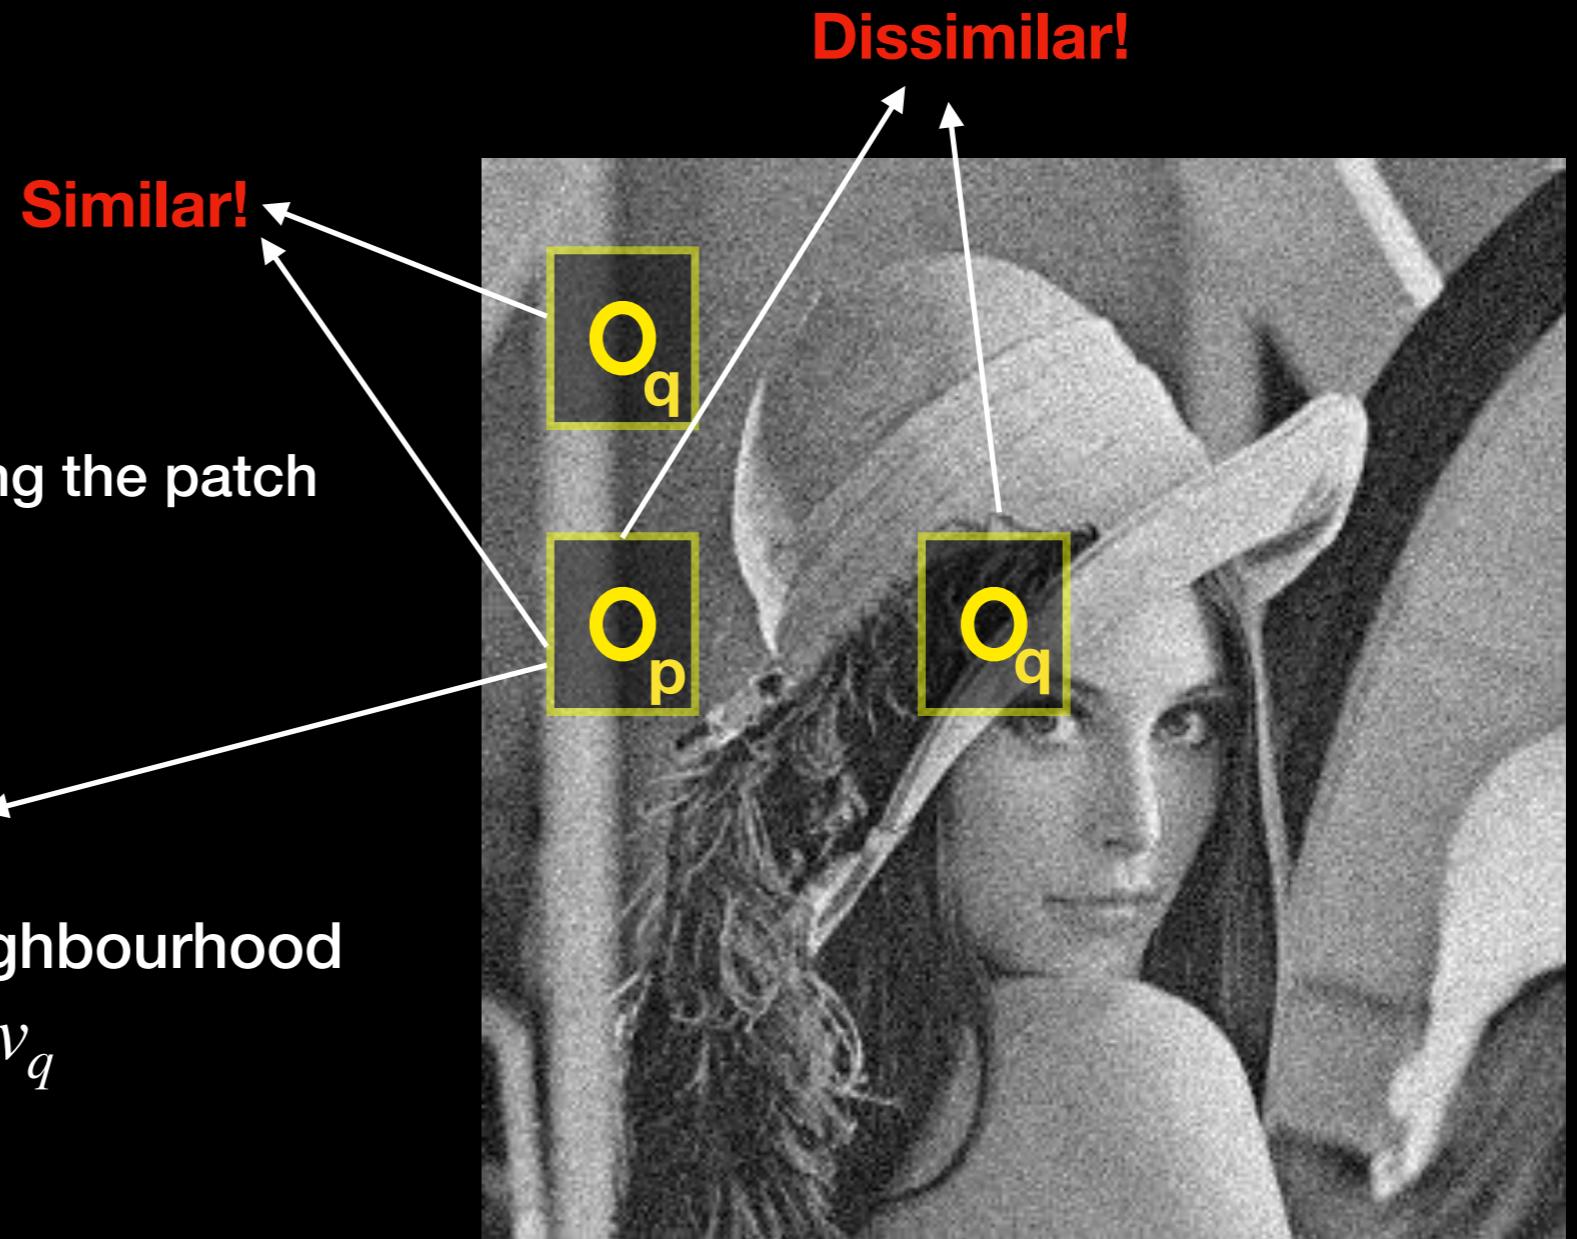
$$\|v_p - v_q\|^2$$

v_p, v_q are vectors describing the patch

for every pixel p

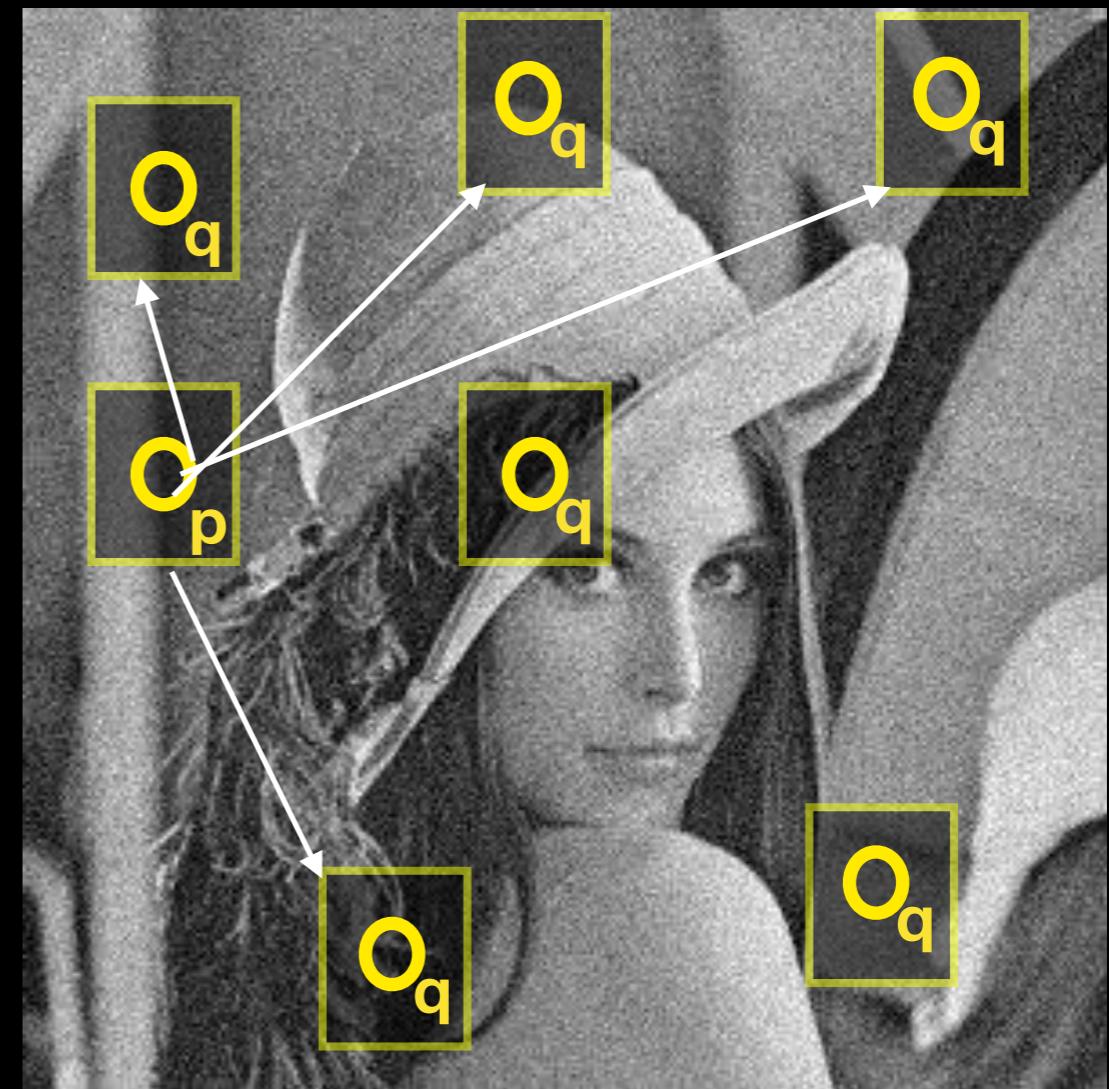
- define a small neighbourhood
- define vectors v_p, v_q

$$NLM[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_r}(\|v_p - v_q\|^2) I_p$$



Non-local means: complexity?

$$NLM[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_r}(\|v_p - v_q\|^2) I_p$$



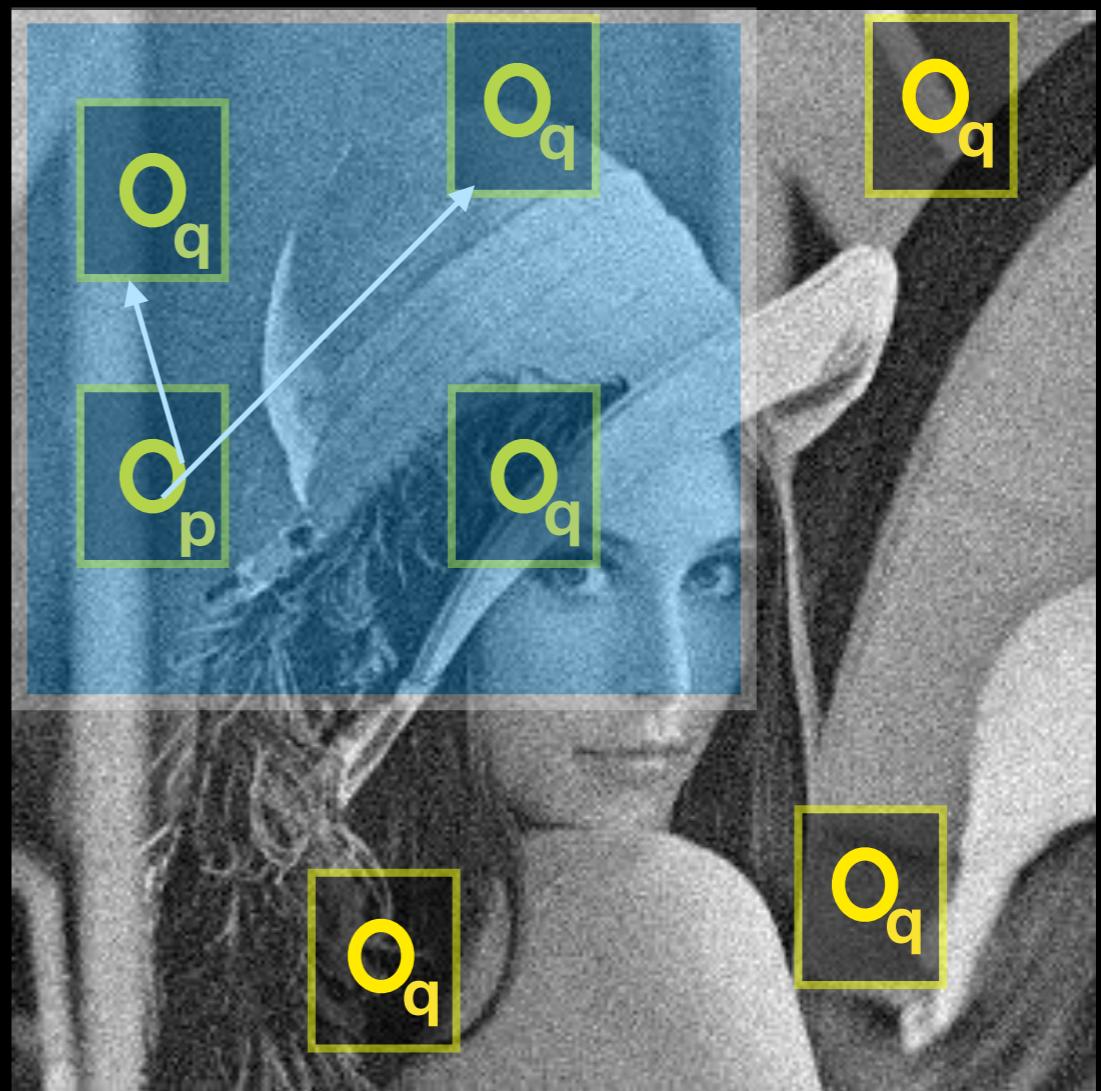
Problem: Too many patches to search!

Non-local means: complexity?

$$NLM[I_p] = \frac{1}{W} \sum_{q \in S} G_{\sigma_r}(\|v_p - v_q\|^2) I_p$$

Problem: Too many patches to search!

Simple remedy: confine the search to an extended neighbourhood.



Non-local means: faster?

PyNotebook

- Get all the patches from the image
- Compute their PCA — know which patterns are more descriptive than others.
- Since low-importance components are mostly noise — we can discard them. We reduce the patch size to a smaller number.
- Index all the reduced patch-vectors in a tree and query over it.